

# CS 6023 - GPU Programming

# Parallel Computing

# Architecture

03/08/2018

# Agenda

How do we contrast CPUs from GPUs?

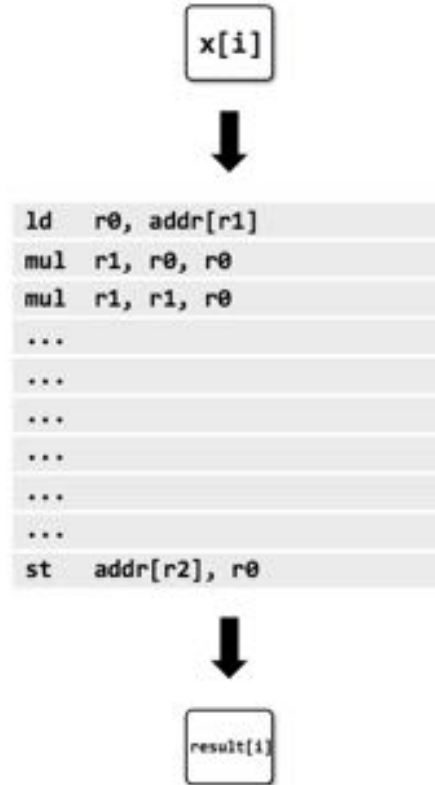
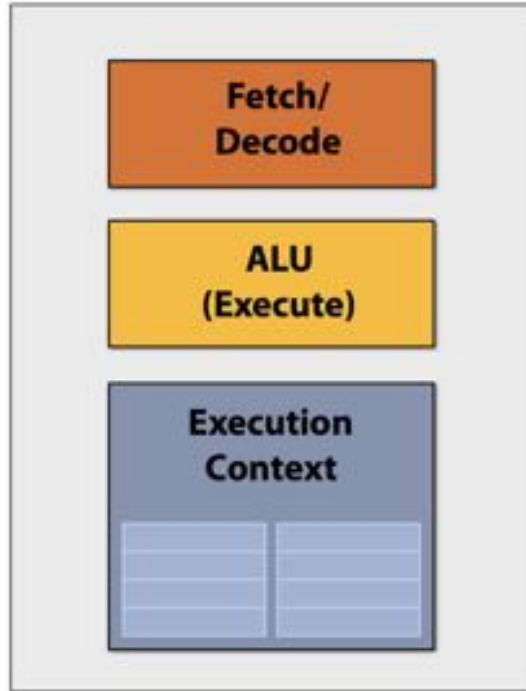
Part 1 of 2:

**Last time:** CPU architecture background

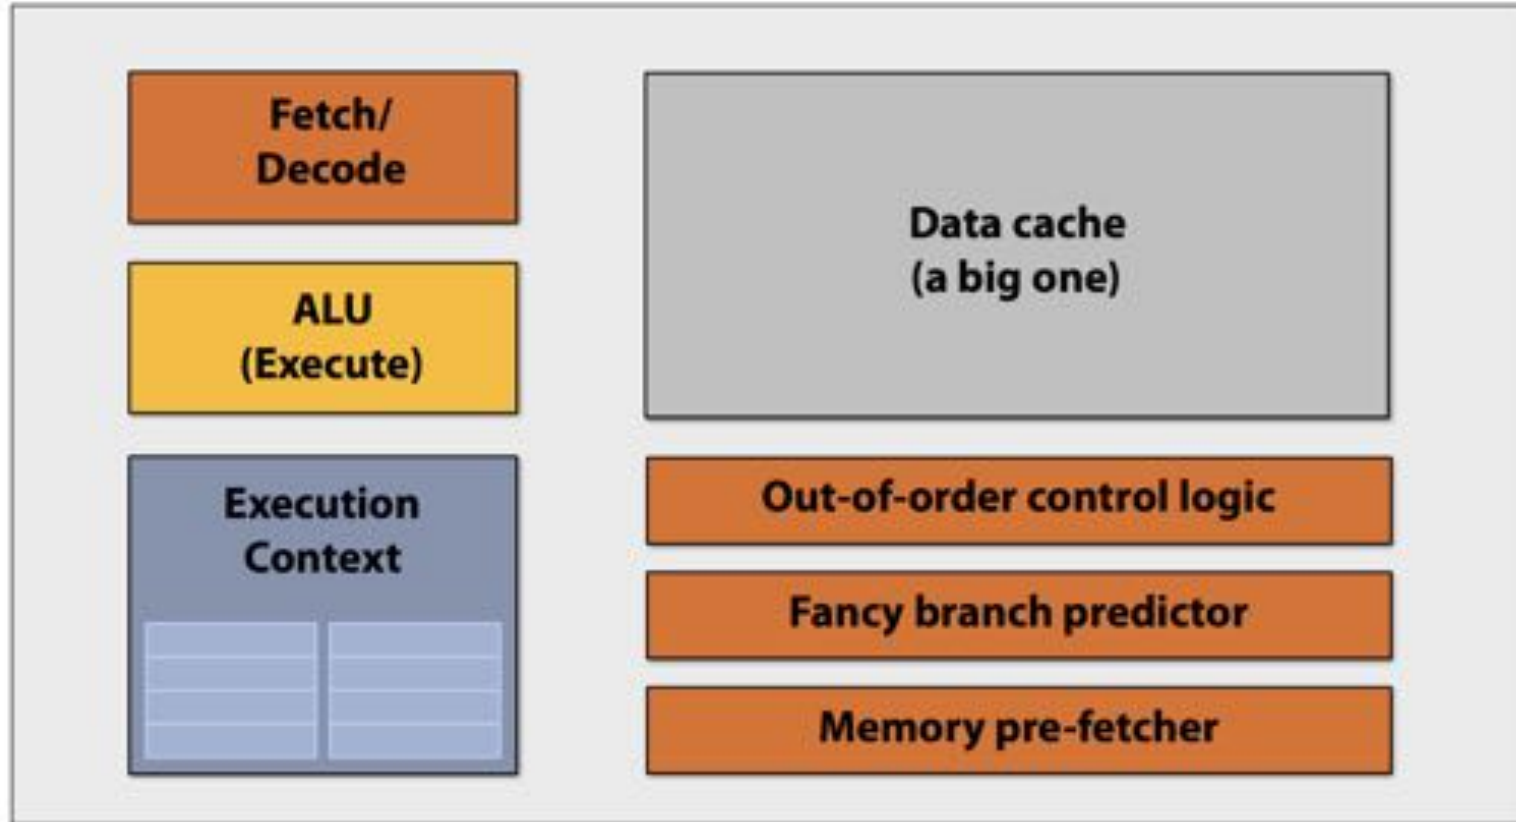
Part 2 of 2:

**Today:** Parallel computer architecture background

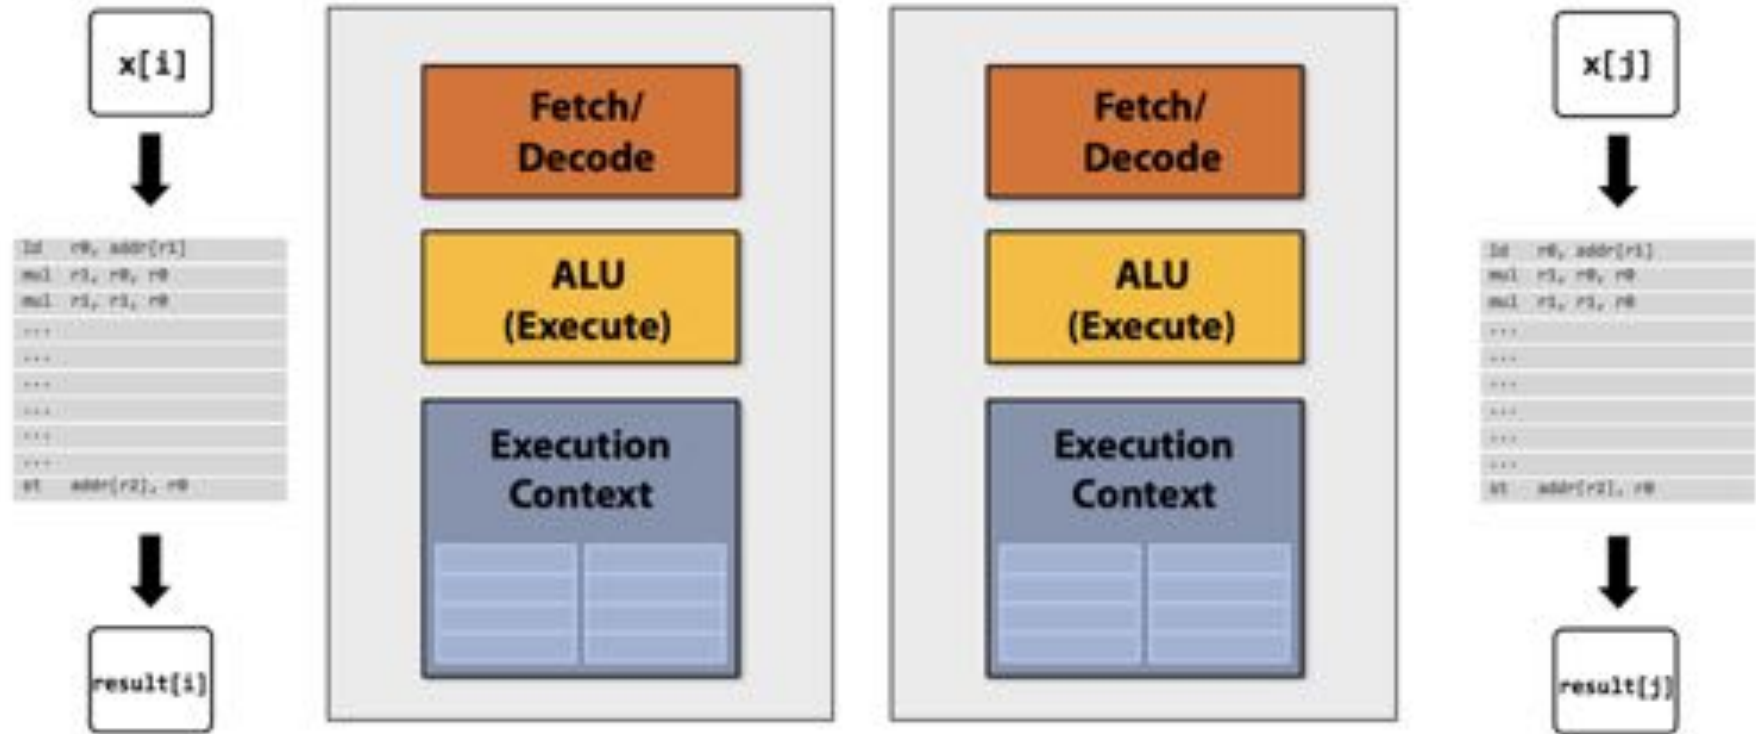
## CPU - single core cartoon



## CPU cartoon - single core era



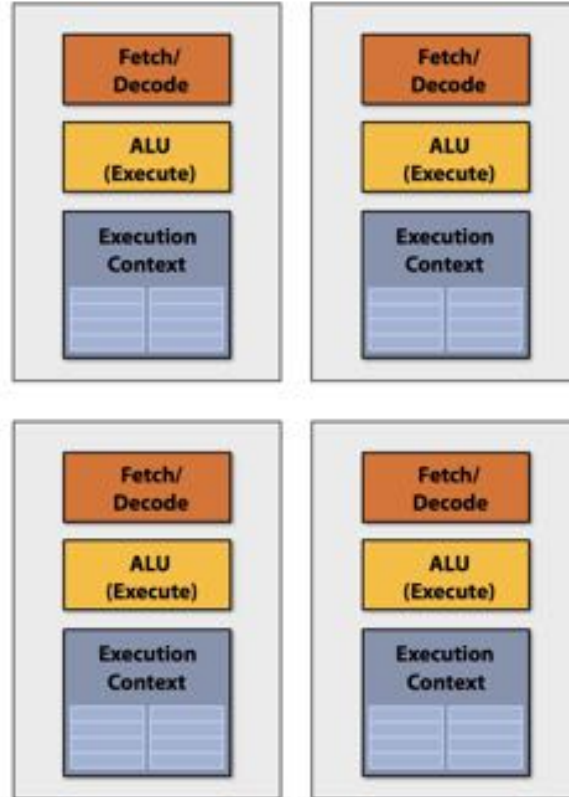

# CPU cartoon - multi-core





Make cores simpler, have multiple of them

(CMU 15-418/618)

# CPU cartoon - many core

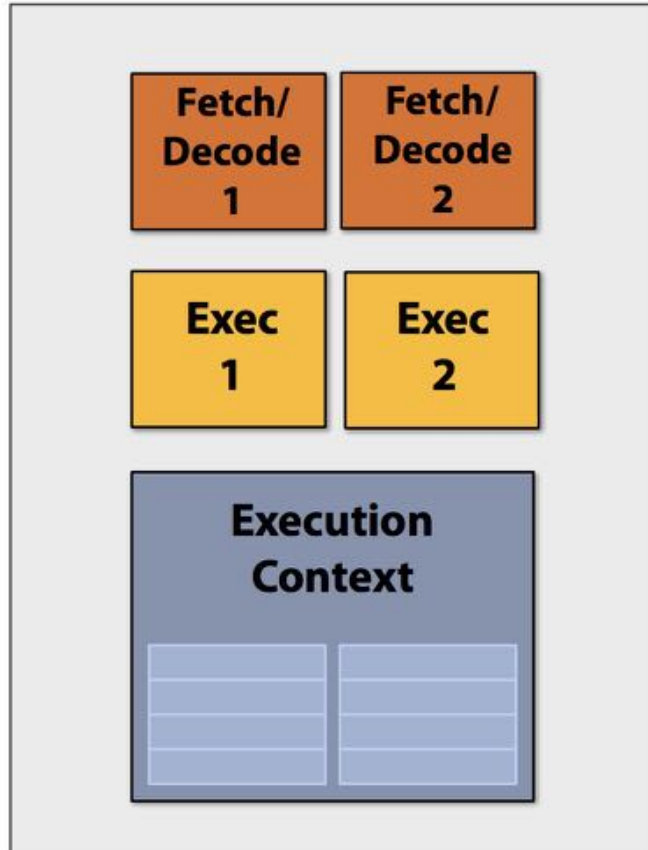


## CPU cartoon - many core



Why don't we  
keep doing  
this for better  
and better  
CPUs?

# What is this?



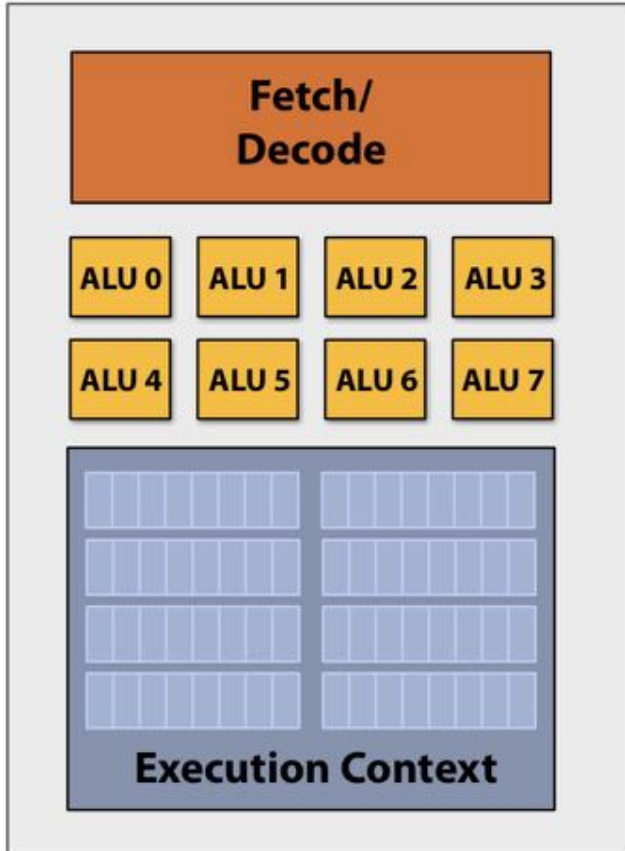
Superscalar

When is this useful?

- For exploiting ILP



# What is this?




SIMD processing

When is this useful?

- Explicit SIMD
- Vector processing

# What is this?





Multithreading

When is this useful

- For hiding latency during stalls

# Combining multicore and SIMD



How many separate instruction streams does this process

16 cores, 16 simultaneous instruction streams

# SIMD processing

```
void sinx(int N, int terms, float* x, float* result)
{
    // declare independent loop iterations
    forall (int i from 0 to N-1)
    {
        float value = x[i];
        float numer = x[i] * x[i] * x[i];
        int denom = 6; // 3!
        int sign = -1;


        for (int j=1; j<=terms; j++)
        {
            value += sign * numer / denom
            numer *= x[i] * x[i];
            denom *= (2*j+2) * (2*j+3);
            sign *= -1;
        }

        result[i] = value;
    }
}
```

Each loop has same set of instructions  
to be applied independently on different  
sets of data

=> SIMD processing

## But we have a problem with conditions



(assume logic below is to be executed for each element in input array 'A', producing output into the array 'result')

```
<unconditional code>


float x = A[i];

if (x > 0) {
    float tmp = exp(x, 5.f);
    tmp *= kMyConst1;
    x = tmp + kMyConst2;
} else {
    float tmp = kMyConst1;
    x = 2.f * tmp;
}

<resume unconditional code>

result[i] = x;
```


# But we have a problem with conditions



# But we have a problem with conditions

Time (clocks)

1 2 ... 8  
ALU 1 ALU 2 ... ALU 8



**Not all ALUs do useful work!**

**Worst case: 1/8 peak performance**

(assume logic below is to be executed for each element in input array 'A'; producing output into the array 'result')

<unconditional code>

```
float x = A[i];
```

```
if (x > 0) {
```

```
    float tmp = exp(x, 5.f);
```

```
    tmp *= kMyConst1;
```

```
    x = tmp + kMyConst2;
```

```
} else {
```

```
    float tmp = kMyConst1;
```

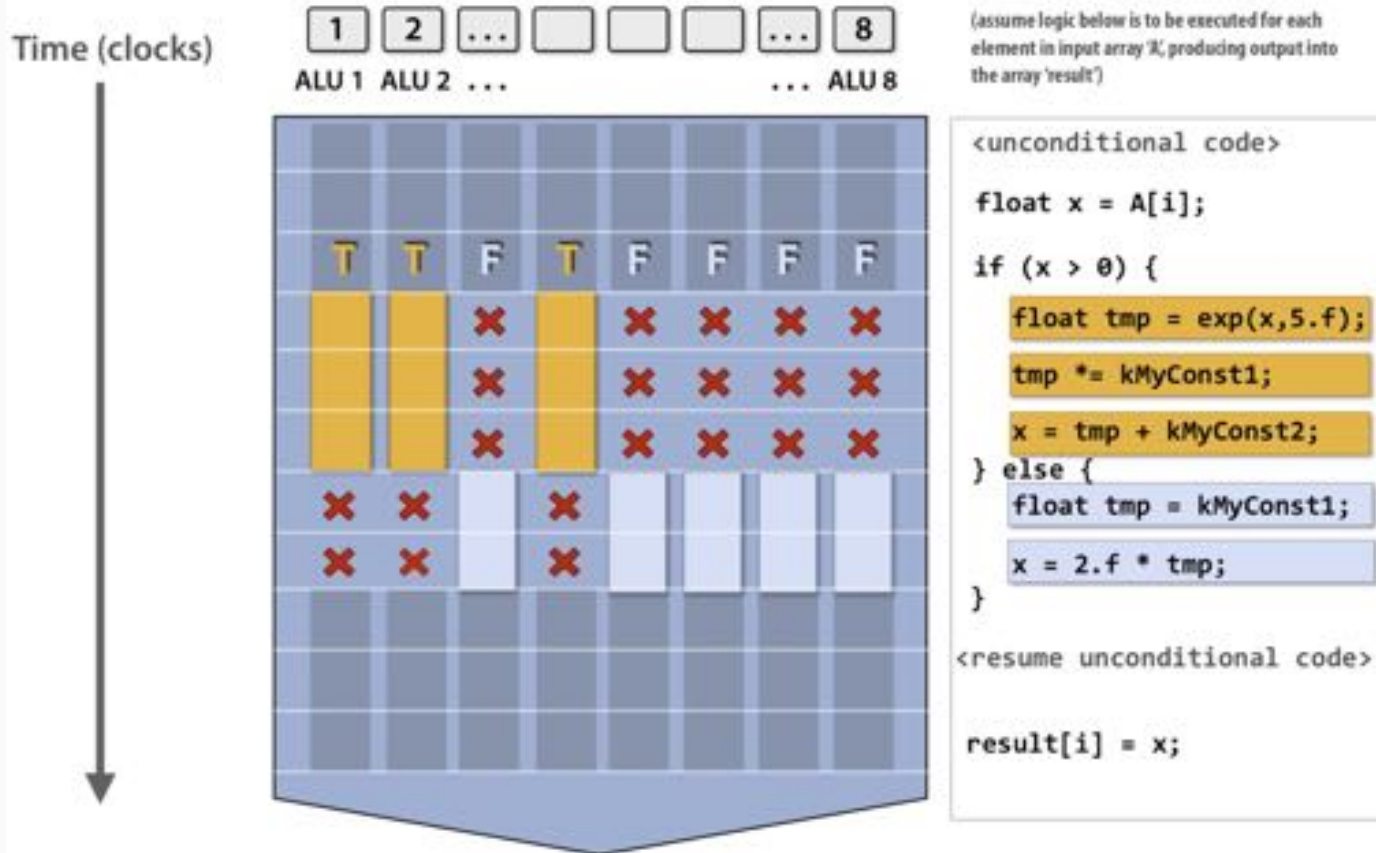
```
    x = 2.f * tmp;
```

```
}
```

<resume unconditional code>

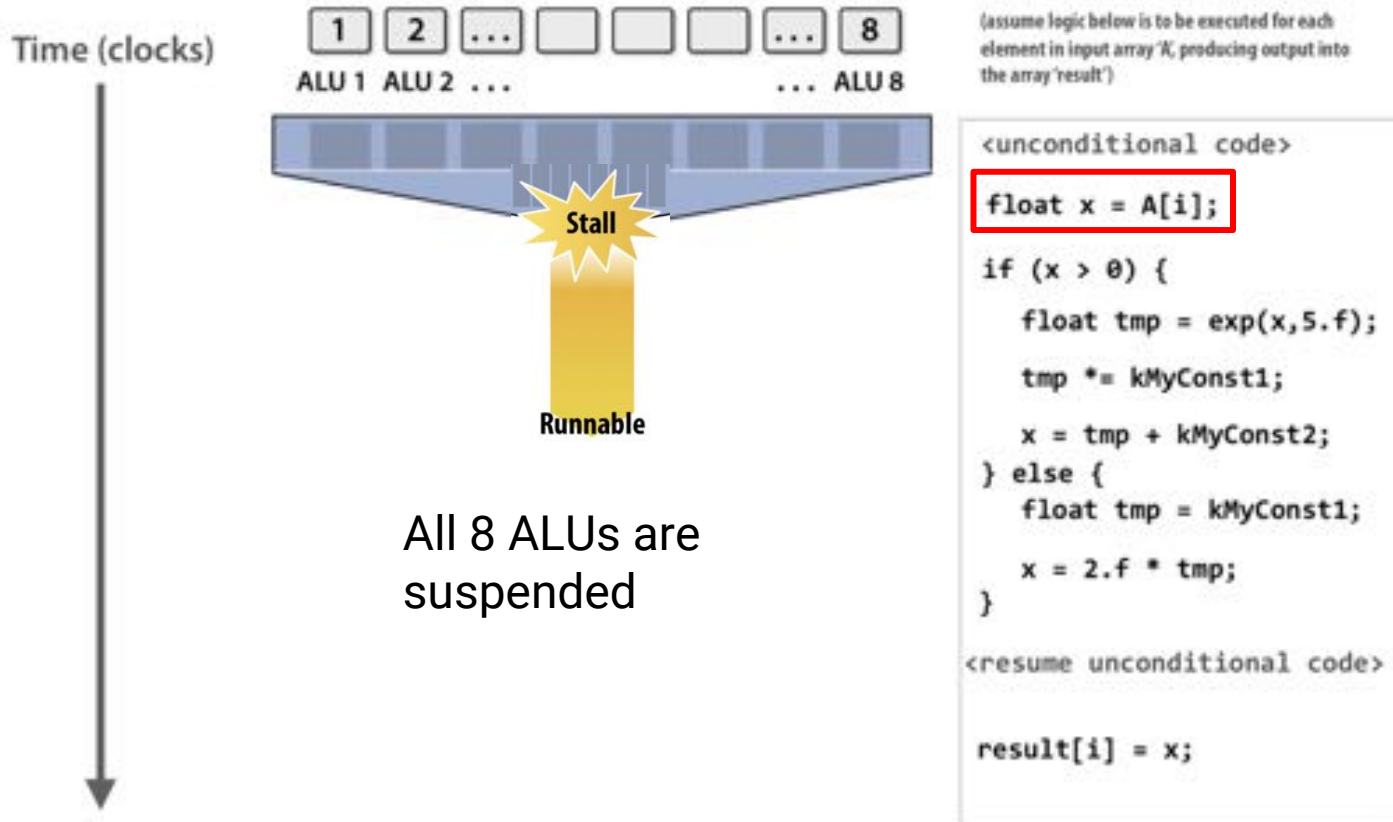
```
result[i] = x;
```

# But we have a problem with conditions



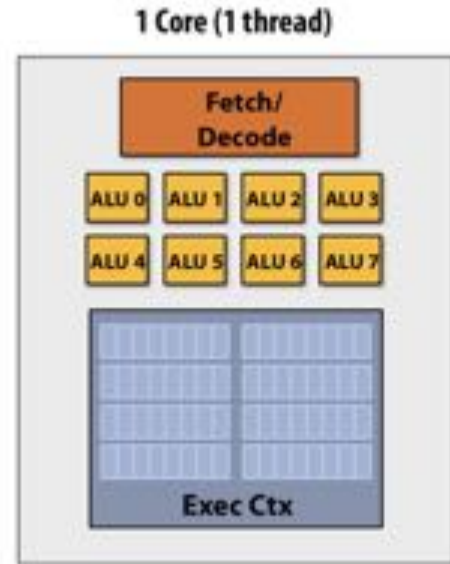


## Problem 2: What if there is a high latency stall

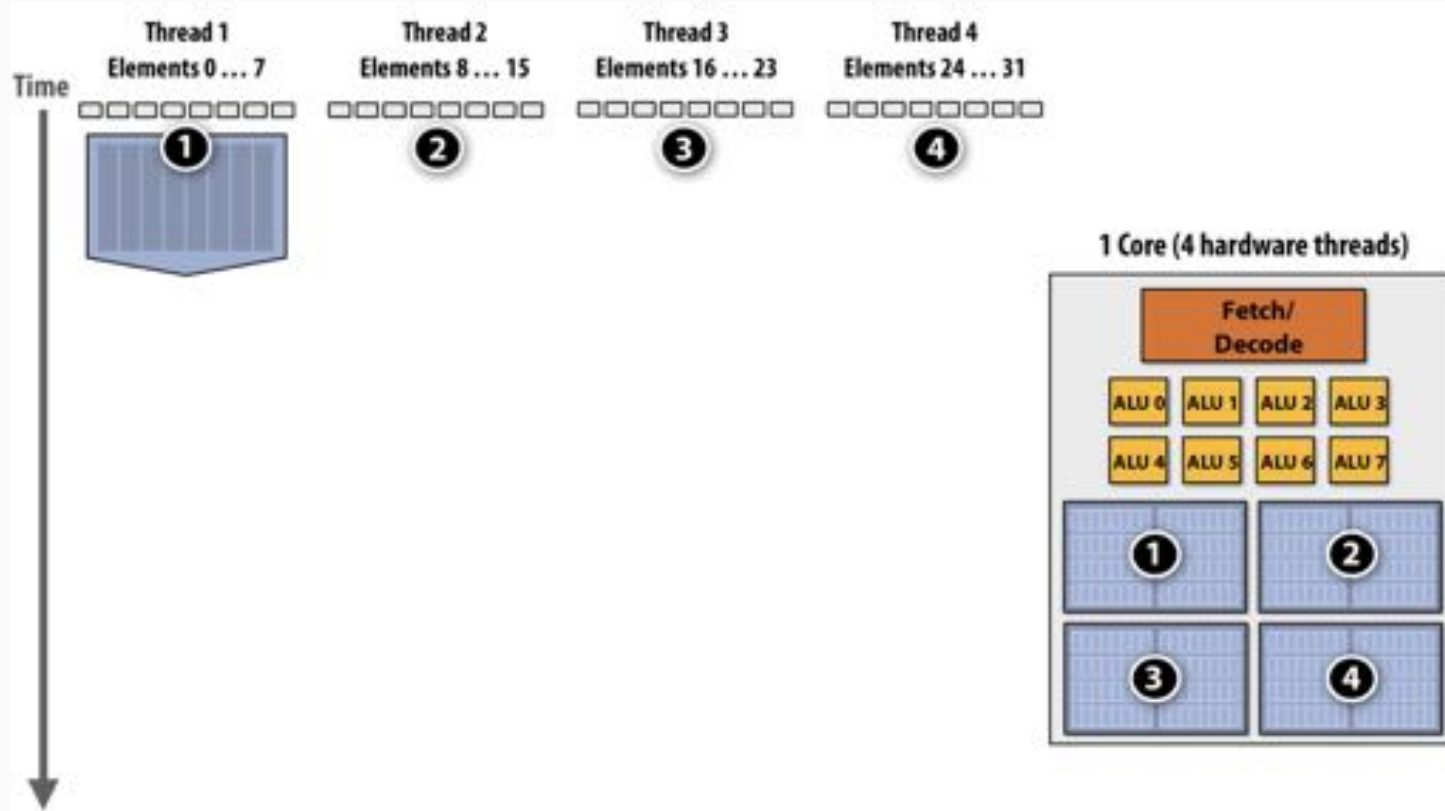


Assume a large latency in memory access

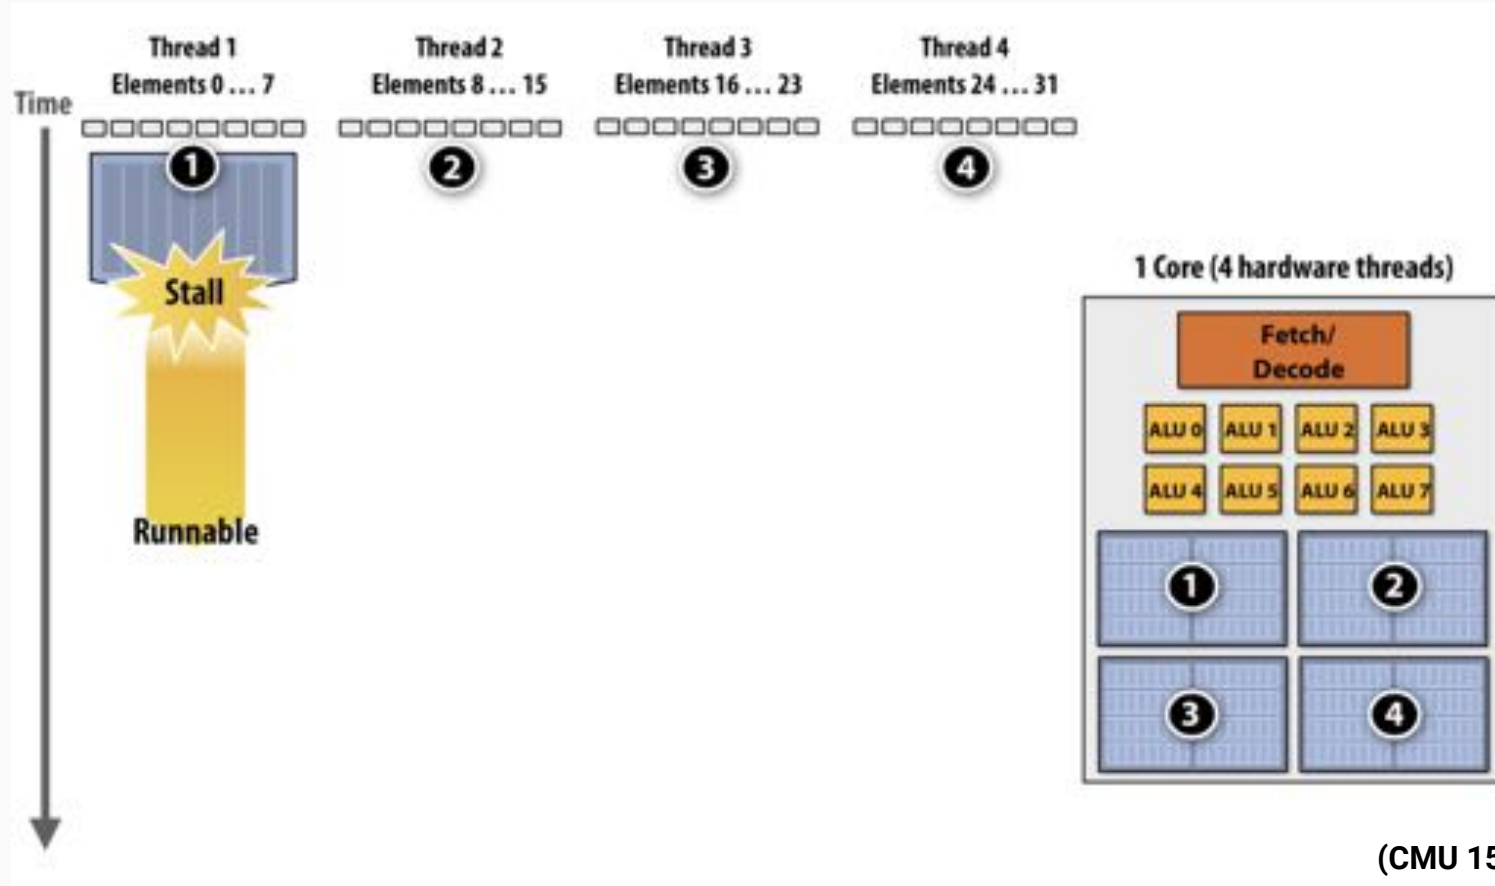
# SIMD + Multithreading



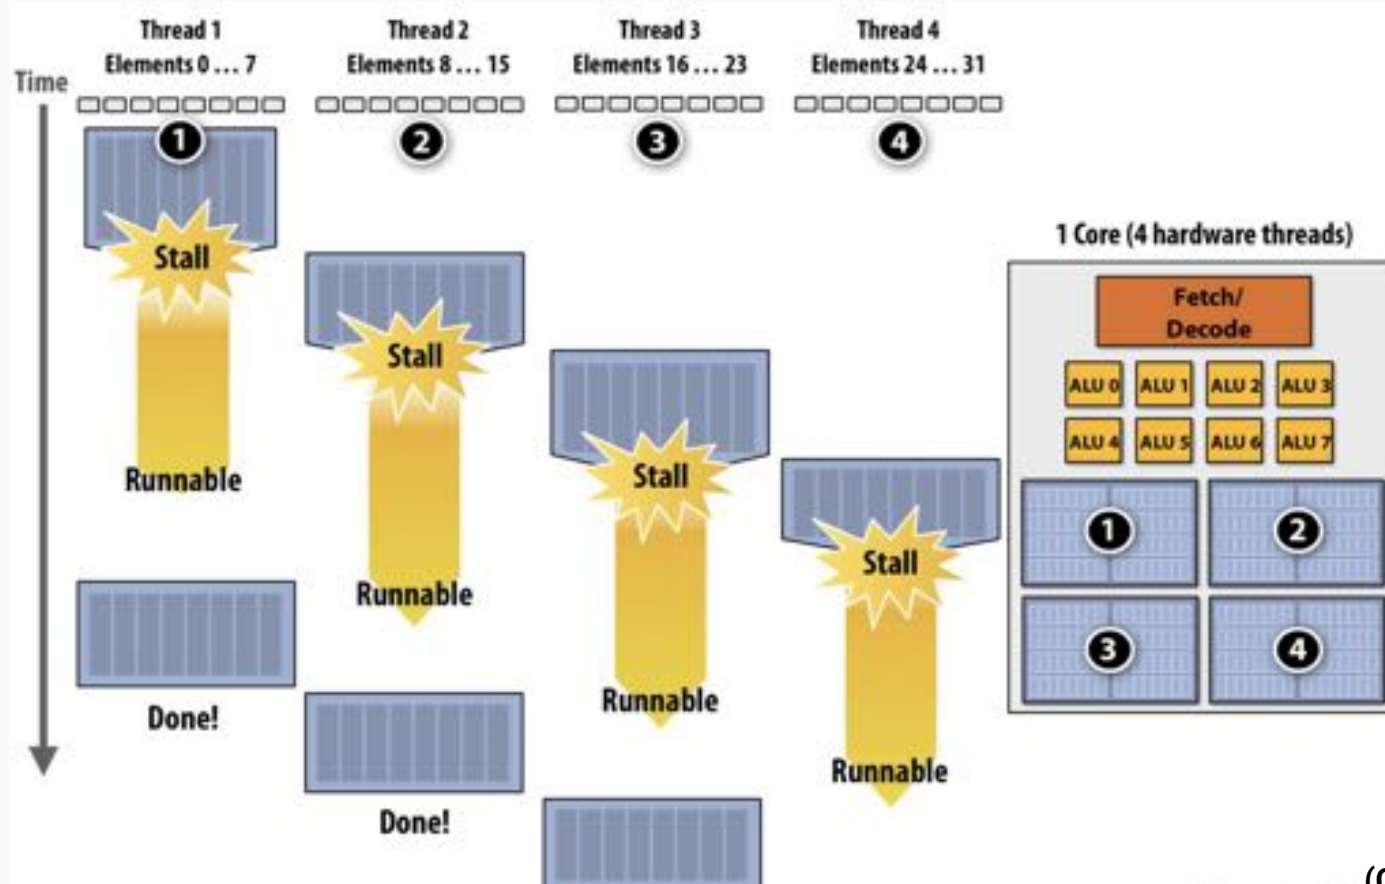
# Partition context to support multiple threads



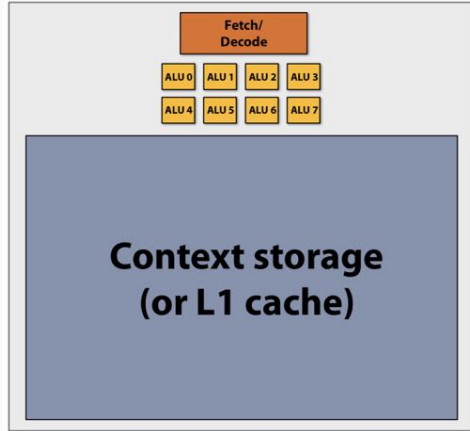
If one thread blocks, run the other



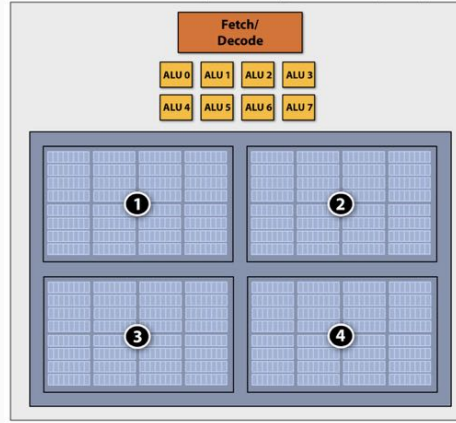
# Have enough blocks to hide latency



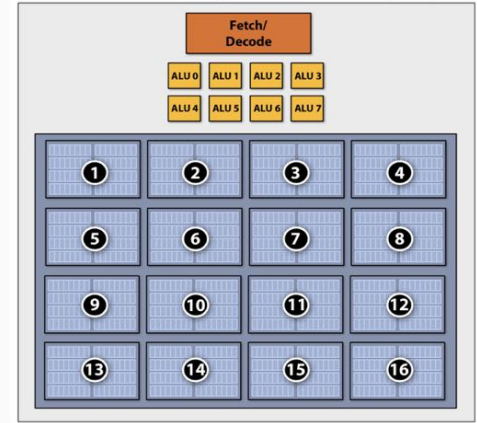
# Trade off in context size




VS



VS



## Fictitious chip



Number of cores **16**

SIMD ALUs per core **8**

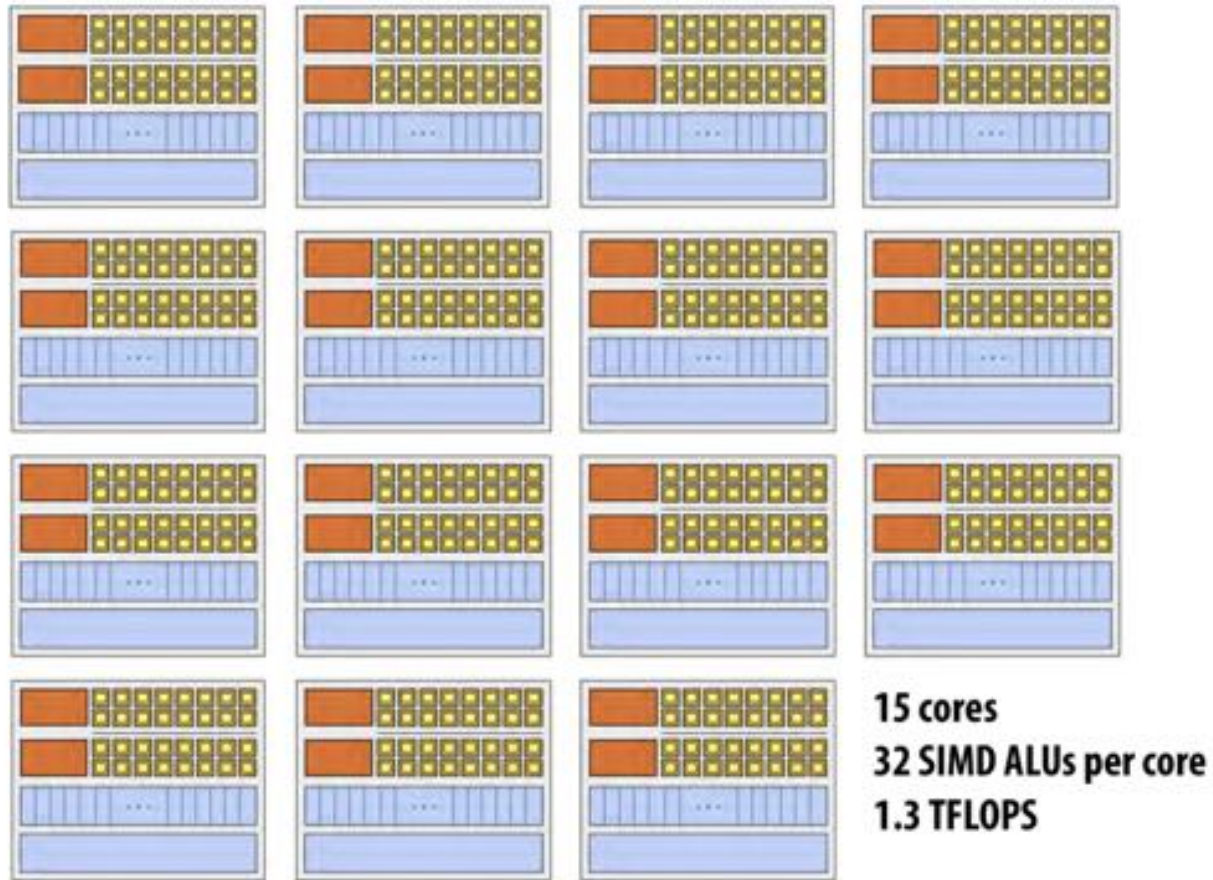
Threads per core **4**

Simultaneous instruction streams **16**

Concurrent instruction streams **64**


Total pieces of work to fully utilize processor **512**

# Nvidia GTX 480





# The GPU architecture




Each core has 16 ALUs running at twice the clock speed => SIMD instruction working on 32 data items

We will call each of these a **thread**  
And a block of 32 threads a **warp**

Up to 48 warps are simultaneously scheduled => 1536 threads in flight

# Roofline model

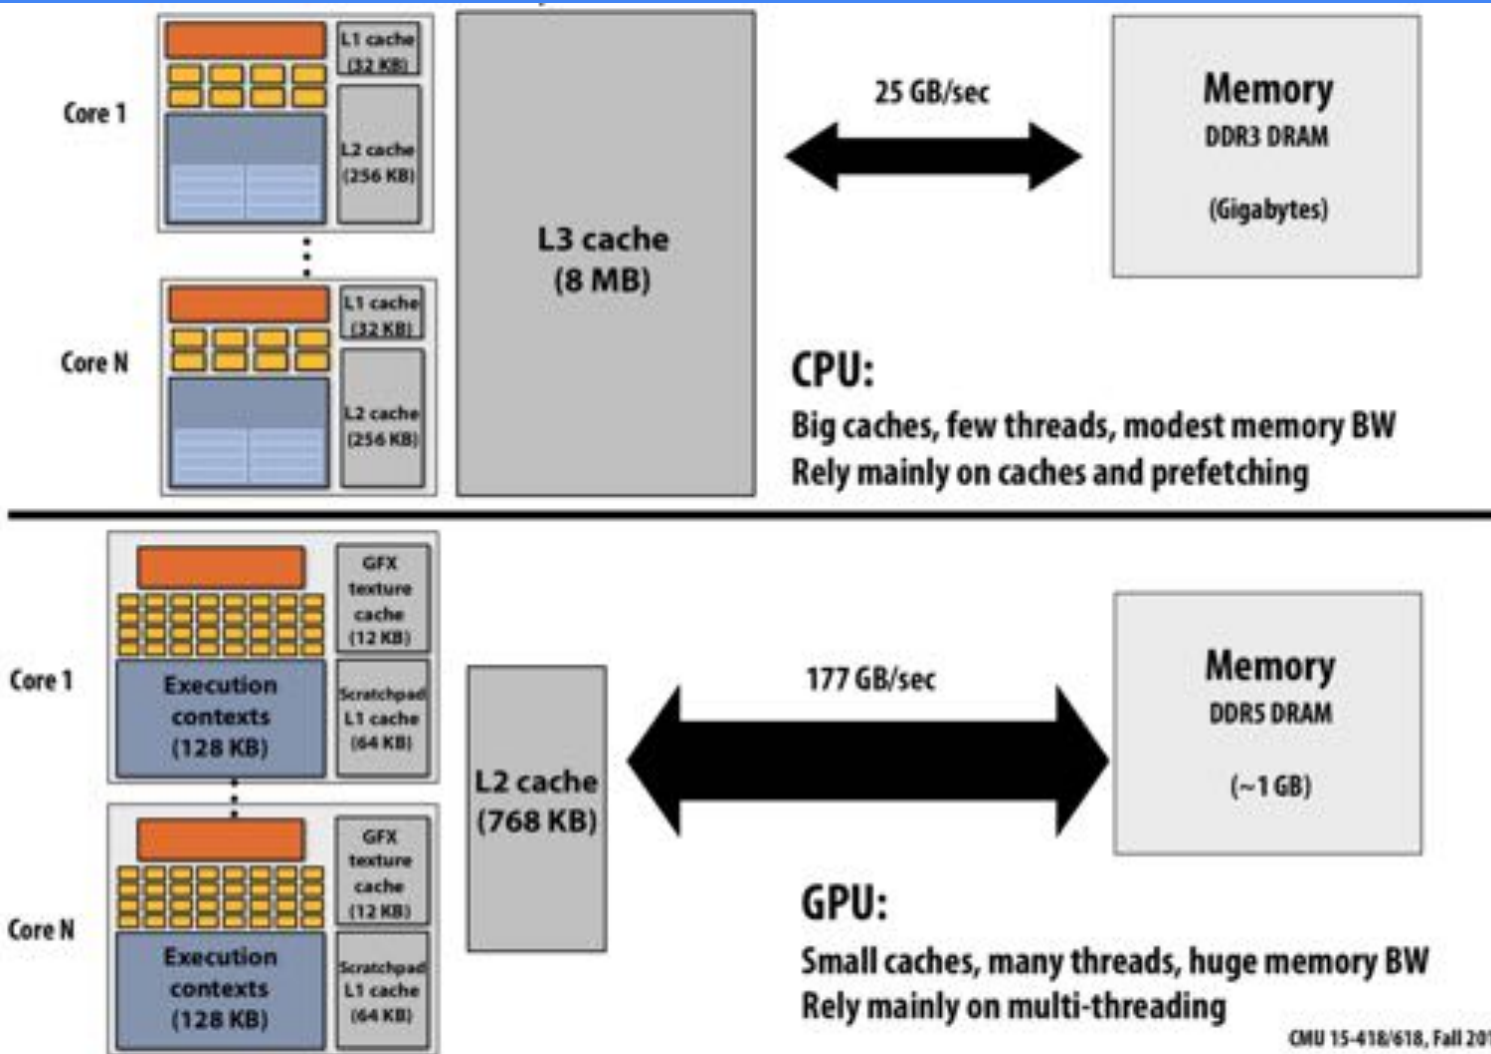


## Exercise

Given a 1 TFLOPS machine, how much memory bandwidth is required to avoid being memory bound for

- vector point-wise multiplication (assume 4 bytes per word)
- [homework] FFT

# Memory



Next time

Start with CUDA programming!