

CS 6023 - GPU Programming

Parallel Reduce Operation

12/10/2018

Setting and Agenda

- We are looking at some common parallel programming patterns and how to optimize them for GPUs
 - Parallel reduce
- Work-efficient and resource-efficient parallel algorithm

Reduction

- For geometric data decomposition, we can divide data into parallel chunks and later **reduce** them to produce the output
- Reduction = Operation that computes a single result from a set of data
- What are some examples from course?

Reduction

- For geometric data decomposition, we can divide data into parallel chunks and later **reduce** them to produce the output
- Reduction = Operation that computes a single result from a set of data
- What are some examples from course?
 - Histogram computation
 - Convolution in input/output stationary data flows
 - More generally, this is common in other programming interfaces (eg. Map-Reduce by Google)

Examples of reduction operation

- Examples
 - Min
 - Max
 - Product
 - Sum
- When does this work?

Examples of reduction operation

- Examples
 - Min
 - Max
 - Product
 - Sum
- When does this work in parallel?
 - Commutative and associative
 - Identity value for initialization

Trivial solutions

- Consider a reduction (say max) on N numbers. What is the complexity of
- Sequential program
- Parallel program

Trivial solutions

- Consider a reduction (say max) on N numbers. What is the complexity of
- Sequential program
 $O(N)$ - Process each value in one step
- Parallel program

Trivial solutions

- Consider a reduction (say max) on N numbers. What is the complexity of
- Sequential program
 $O(N)$ - Process each value in one step
- Parallel program
What does complexity mean?


Parallel solutions

- Simple approach - Apply reduction operation recursively to pairs
 - $O(\log N)$ time
 - $O(N)$ operations
 - $O(N)$ processors

Parallel solutions

- Simple approach - Apply reduction operation recursively to pairs
 - $O(\log N)$ time
 - $O(N)$ operations
 - $O(N)$ processors
- Speed-up vs sequential = $(N - 1) / \log N$
- If N is very large, then parallelism is large, but number of processors required is also large

Parallel reduce - Simple code



- In-place modification of data
- Number of threads gets halved in each iteration
- Maximum number of threads is $N/2$

```
for i = 0 to  $\log_2 N - 1$   
  for all j = 0 to N - 1 by  $2^{i+1}$  in parallel  
     $x[j + 2^{i+1} - 1] += x[j + 2^i - 1];$ 
```

- Usual questions: What will each thread/block do?

- Usual questions: What will each thread/block do?
- Let a block of N threads read in $2*N$ values

```
__shared__ float partialSum[2*BLOCK_SIZE];

unsigned int t = threadIdx.x;
unsigned int start = 2*blockIdx.x*blockDim.x;

partialSum[t] = input[start + t];
partialSum[blockDim.x+t] = input[start + blockDim.x+t];

for (unsigned int stride = 1; stride <= blockDim.x; stride *= 2) {
    if (t % stride == 0)
        partialSum[2*t] += partialSum[2*t+stride];
}
```

- Usual questions: What will each thread/block do?
- Let a block of N threads read in $2*N$ values

```
__shared__ float partialSum[2*BLOCK_SIZE];

unsigned int t = threadIdx.x;
unsigned int start = 2*blockIdx.x*blockDim.x;

partialSum[t] = input[start + t];
partialSum[blockDim+t] = input[start + blockDim.x+t];

for (unsigned int stride = 1; stride <= blockDim.x; stride *= 2) {
    if (t % stride == 0)
        partialSum[2*t] += partialSum[2*t+stride];
}
```

Are we missing something?

- Usual questions: What will each thread/block do?
- Let a block of N threads read in $2*N$ values

```
__shared__ float partialSum[2*BLOCK_SIZE];

unsigned int t = threadIdx.x;
unsigned int start = 2*blockIdx.x*blockDim.x;

partialSum[t] = input[start + t];
partialSum[blockDim.x+t] = input[start + blockDim.x+t];

for (unsigned int stride = 1; stride <= blockDim.x; stride *= 2) {
    __syncthreads();
    if (t % stride == 0)
        partialSum[2*t] += partialSum[2*t+stride];
}
```


- Have we solved the problem?

- Have we solved the problem?
- We have summed up only the values assigned to multiple threads of a block? What about the global sum?

- Do you see any performance bugs?

```
__shared__ float partialSum[2*BLOCK_SIZE];

unsigned int t = threadIdx.x;
unsigned int start = 2*blockIdx.x*blockDim.x;

partialSum[t] = input[start + t];
partialSum[blockDim+t] = input[start + blockDim.x+t];

for (unsigned int stride = 1; stride <= blockDim.x; stride *= 2) {
    __syncthreads();
    if (t % stride == 0)
        partialSum[2*t] += partialSum[2*t+stride];
}
```

- Do you see any performance bugs?

```
__shared__ float partialSum[2*BLOCK_SIZE];

unsigned int t = threadIdx.x;
unsigned int start = 2*blockIdx.x*blockDim.x;


partialSum[t] = input[start + t];
partialSum[blockDim+t] = input[start + blockDim.x+t];

for (unsigned int stride = 1; stride <= blockDim.x; stride *= 2) {
    __syncthreads();
    if (t % stride == 0)
        partialSum[2*t] += partialSum[2*t+stride];
}
```

Lot of control
divergence on
adjacent warps

Optimization

- Remove control divergence amongst consecutive warps by compacting
- Exercise: Try out the CUDA kernel



- Class exercise (solution added at the end)

Exercise

- For a 1024 thread block, what are the number of warps with control divergence in each iteration

Exercise

- For a 1024 thread block, what are the number of warps with control divergence in each iteration
 - First 5 iterations (1024, 512, 256, 128, 64, 32 threads with work) there is NO control divergence
 - For the next 5 iterations 1 warp will have control divergence

Work and resource efficiency

- What is the total work done by the sequential algorithm? $N-1$ adds
- What is the total work done by the parallel algorithm? $N-1$ adds
- When both sequential and parallel algorithms do the same work, we say the parallel algorithm is **work-efficient**
- Many parallel algorithms are not work efficient (eg. prefix sum in next class)

Work and resource efficiency

- What is the total work done by the sequential algorithm? $N-1$ adds
- What is the total work done by the parallel algorithm? $N-1$ adds
- When both sequential and parallel algorithms do the same work, we say the parallel algorithm is **work-efficient**
- Many parallel algorithms are not work efficient (eg. prefix sum in next class)

- But the parallel reduce algorithm is not resource efficient
- In sequential case: 1 thread runs for $N-1$ time
- In parallel case: $N/2$ threads are reserved at the start for $\log(N)$ time
- In some algos, trade-off exists between resource usage and work efficiency

CUDA code for compact parallel reduce with lesser control divergence

```
__shared__ float partialSumPing[2*BLOCK_SIZE], partialSumPong[2*BLOCK_SIZE];

unsigned int t = threadIdx.x;
unsigned int start = 2*blockIdx.x*blockDim.x;

partialSumPing[t] = input[start + t];
partialSumPing[blockDim+t] = input[start + blockDim.x+t];
partialSumPong[t] = 0;
partialSumPong[blockDim+t] = 0;

bool odd = true;

for (unsigned int stride = blockDim.x; stride > 0; stride /= 2) {
    __syncthreads();
    if (t < stride) {
        if (odd)
            partialSumPong[t] += partialSumPing[t+stride];
        else
            partialSumPing[t] += partialSumPong[t+stride];
        odd = !odd;
    }
}
```

Why do we need two arrays?