

CS 6023 - GPU Programming

Parallel Scan Operation

17/10/2018

Quiz

- Scheduled on Mon 22nd
- Interchange of A and G slots => **Quiz at 8am**
- Different format this time
 - Subjective questions
 - Write code, give explanations, derive formulae etc.
 - Syllabus: Lectures 7, 8, 9, 10, 11 (today)
 - Marks: $3 \times 3 + 1$ (for neatness/concise answers)
 - Allowed one sheet of paper with notes

Setting and Agenda

- We are looking at some common parallel programming patterns and how to optimize them for GPUs
 - Parallel scan
- Work-efficiency in parallel algorithms

- Given a binary associative \otimes operation and an array of n elements

$$[x_0, x_1, \dots, x_{n-1}]$$

the **scan operation** computes the array

$$[x_0, x_0 \otimes x_1, \dots, x_0 \otimes x_1 \otimes \dots x_n]$$

- Given a binary associative \otimes operation and an array of n elements

$$[x_0, x_1, \dots, x_{n-1}]$$

the **scan operation** computes the array

$$[x_0, x_0 \otimes x_1, \dots, x_0 \otimes x_1 \otimes \dots \otimes x_n]$$

- When the operation is addition, the scan operation computes the prefix sum or the cumulative sum
- Example: $[5\ 2\ 1\ 3\ 6\ 7\ 0\ 4] \rightarrow [5\ 7\ 8\ 11\ 17\ 24\ 24\ 28]$
- It has a wide set of applications in many algorithms from sorting, string comparisons, to statistics

Exclusive and Inclusive Scans

- Inclusive scan
 - X: [5 2 1 3 6 7 0 4]
 - Y: [5 7 8 11 17 24 24 28]
- Exclusive scan
 - X: [5 2 1 3 6 7 0 4]
 - Y: [0 5 7 8 11 17 24 24]
- Exercise: Given a kernel for inclusive scan, how do you compute exclusive scan? Vice versa?

Very simple sequential algorithm

Very simple sequential algorithm

```
y[0] = x[0];  
for (int k = 1; k < N; ++k)  
    y[k] = y[k - 1] + x[k];
```

- Complexity: $O(N)$ operations
- For sequential programs: Parallel reduce and scan have almost same amount of work
- We have a loop (like in matrix multiplication / convolution / parallel reduce). Can we use loop-parallelism here?

The simplest parallel program

The simplest parallel program


In thread i , compute $y[i]$ as
$$y[i] = x[0] + \dots + x[i]$$

The simplest parallel program


In thread i , compute $y[i]$ as
 $y[i] = x[0] + \dots + x[i]$

- Complexity:
 - $O(N^2)$ operations
 - $O(N)$ time
 - $O(N)$ threads


Next best solution




Next best solution




Next best solution




Next best solution




Next best solution




Next best solution



Next best solution




Next best solution




Next best solution

End of
iteration 1
running in
parallel



Do we have a
problem with
in-place edit?


Next best solution



Iteration 2

Next best solution

Iteration 3



Prove that
this process
computes the
prefix sum

What is the complexity?


What is the complexity?

- Complexity:
 - $O(\log(N))$ time
 - $O(N)$ resources
 - Operations?

What is the complexity?


- Complexity:
 - $O(\log(N))$ time
 - $O(N)$ resources
 - Operations?
 - Iteration i does $(N - 2^i)$ operations
 - $\log(N)$ number of iterations
 - Total operations: $N \cdot \log(N) - (N - 1) \Rightarrow O(N \cdot \log(N))$
- $O(N \cdot \log(N))$ is an improvement, but it is clearly not work efficient

Can we do even better?




- Consider parallel reduce:
At the end some indices have the correct prefix sum
- It has an efficient $O(N)$ complexity of operations
- Formally, what is the value at each index?


Second phase



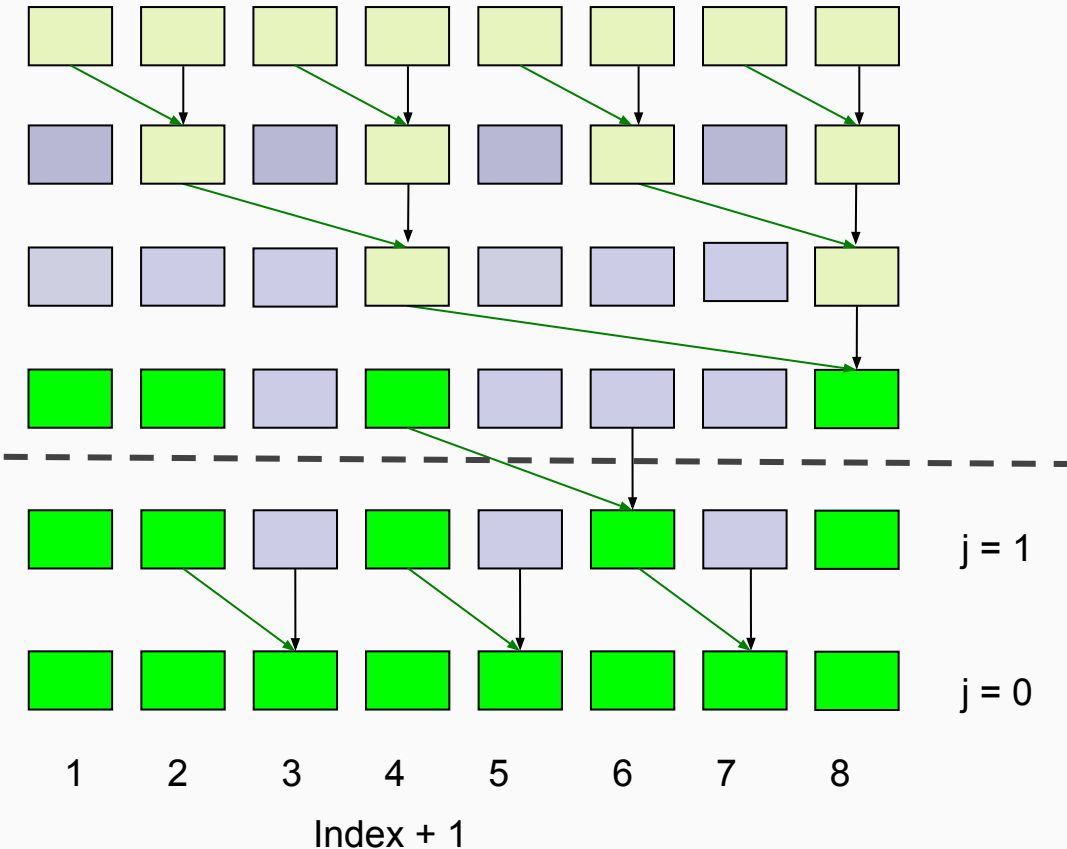
Second phase



Second phase




Second phase




- Need to update only certain values in downward pass
- Which ones?
- Value at (index+1) is updated at step j , if its largest factor which is a power of 2 is 2^j and it is not a power of 2

Consider $N = 16$




Phase 1: parallel reduce

Consider $N = 16$




Phase 2


Consider $N = 16$




Consider $N = 16$



Consider $N = 16$



Consider $N = 16$



Elegant CUDA code for second phase

```
for (unsigned int stride = N/4; stride > 0; stride /= 2) {  
    __syncthreads();  
    int t = (threadIdx.x + 1) * stride * 2 - 1;  
    if (t + stride < N) {  
        partialSum[t + stride] += partialSum[t];  
    }  
}
```

- Complexity:
 - $O(\log(N))$ time
 - $O(N)$ resources
 - Operations?

- Complexity:
 - $O(\log(N))$ time
 - $O(N)$ resources
 - Operations
 - Phase 1 parallel reduce is $O(N)$
 - Phase 2 has $\log(N)-1$ steps
 - Iterations do $(2 - 1)$, $(4 - 1)$, ... $(N/2 - 1)$ operations
 - Total = $(N - \log(N) + 1) \rightarrow O(N)$ operations
 - Total $O(N)$ operations

Is it work efficient

- Recall sequential code had $O(N)$ complexity for scan
- Now, we have $O(N)$ for parallel scan
- But not quite work efficient
 - Phase 1 and 2 can do at most $2x(N - 1)$ adds
 - Thus, upto twice the number of operations as sequential
- But constant factors do not matter because we are interested in asymptotic trend for highly parallel implementations