

Assignment 2 : CS6023 GPU Programming

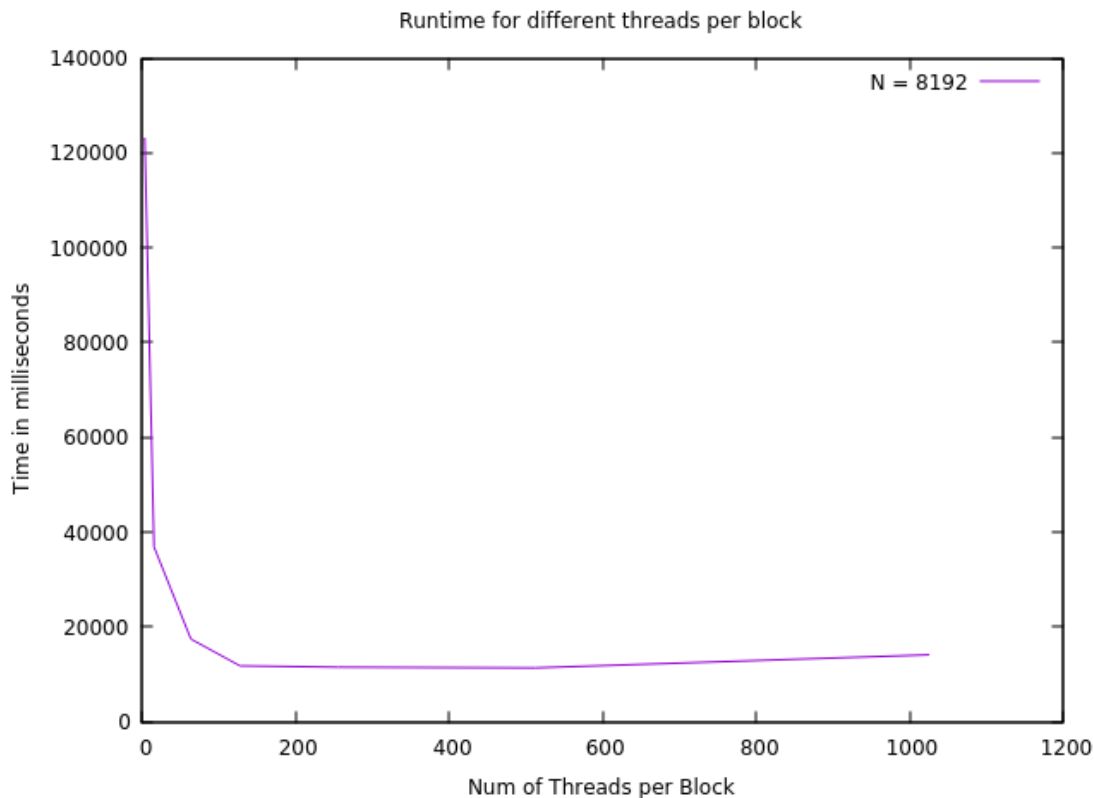
by CS15B049

Q1 : The code is included.

Q2: Matrix Size : 8192*8192

S.No	Threads per Block	Runtime in milliseconds
1	2*2	123207.99
2	4*4	36875.53
3	8*8	17413.35
4	16*8	11756.33
5	16*16	11465.96
6	16*32	11328.14
7	32*32	14089.15
8	64*64	Error as numof threads >1024

The Runtime graph :

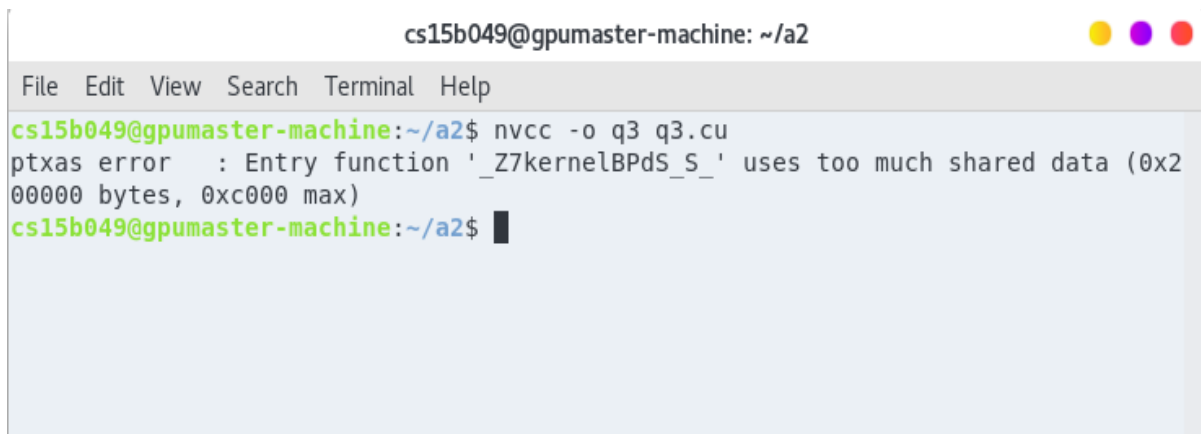


As we increase the number of threads per block the run time decreases exponentially and then after 512(32*16) it increases.

Q3: - Using 16x16 threads per block.

a) The code is included.

b) If we change the size of the array to 8192, the code doesn't compile through nvcc, it is due to the fact that size of shared memory is limited and the size of the array we initialize goes beyond its size. It gives compile time error as arrays are statically allocated and compiler knows the size of shared memory and does a check for the allocation is possible or not.

A screenshot of a terminal window titled 'cs15b049@gpumaster-machine: ~/a2'. The terminal shows the command 'nvcc -o q3 q3.cu' being executed. The output is a 'ptxas error' message: 'Entry function '_Z7kernelBPdS_S_' uses too much shared data (0x200000 bytes, 0xc000 max)'. The prompt returns to the user.

```
cs15b049@gpumaster-machine: ~/a2
File Edit View Search Terminal Help
cs15b049@gpumaster-machine:~/a2$ nvcc -o q3 q3.cu
ptxas error  : Entry function '_Z7kernelBPdS_S_' uses too much shared data (0x200000 bytes, 0xc000 max)
cs15b049@gpumaster-machine:~/a2$
```

Q4: Size of matrix = 8192×8192
Threads per block = 16×16

Runtime using tiling : 3833.963135 milliseconds

Runtime using q1 (without tiling and shared memory) : 11465.964844 milliseconds

as the above with tiling runtime is reduced to 1/3 of without tiling.


It is due to the fact that in tiling we are utilizing shared memory which is much faster than the DRAM global memory, so the latency of accessing the element from DRAM is reduced as we utilize the temporal and spatial locality as well as burst section with the use of tiling, so overall latency reduces.

Q5: Matrix Size : 8192×8192

Number of threads per block = $\text{TILE_SIZE} \times \text{TILE_SIZE}$

Optimal Tile Size = 32×32

below is th runtime graph with diffenrent tile size.



Q6 : Matrix A size : 4096×8192
Matrix B size : 8192×16384

Runtime : 11539.041016 milliseconds.