



CSCI-GA.3033-004

Graphics Processing Units (GPUs): Architecture and Programming

CUDA

Advanced Techniques 2

Mohamed Zahran (aka Z)

mzahran@cs.nyu.edu

<http://www.mzahran.com>




Alignment

Memory Alignment

- Memory access on the GPU works much better if the data items are aligned at 64 byte boundaries.
- Hence, allocating 2D (or 3D) arrays so that every row starts at a 64-byte boundary address will improve performance.
- Difficult to do for a programmer!

Pitch



2D Arrays

- CUDA offers special versions of:
 - Memory allocation of 2D arrays so that every row is padded (if necessary). The function determines the best pitch and returns it to the program. The function name is `cudaMallocPitch()`
 - Memory copy operations that take into account the pitch that was chosen by the memory allocation operation. The function name is `cudaMemcpy2D()`

```
cudaMallocPitch( void** devPtr, ← device array  
                 size_t* pitch, ← Will return the pitch  
                 size_t widthInBytes,  
                 size_t height)
```

- This allocates at least *width (in bytes) X height* array.
- The value returned in pitch is the width in bytes of the allocation.
- The above function determines the best pitch and returns it to the program.
- It is strongly recommends the usage of this function for allocating 2D (and 3D) arrays.

```
cudaError_t cudaMemcpy2D ( void * dst,  
                           size_t dpitch, ← the widths in memory in bytes  
                           const void * src, including any padding added  
                           size_t spitch, ← to the end of each row  
                           size_t width,  
                           size_t height,  
                           enum cudaMemcpyKind kind )
```

- *dst* - Destination memory address
- *dpitch* - Pitch of destination memory
- *src* - Source memory address
- *spitch* - Pitch of source memory
- *width* - Width of matrix transfer (in bytes)
- *height* - Height of matrix transfer (rows)
- *kind* - Type of transfer

Example

```
int main(int argc, char * argv[])
{
    float * A, *dA;
    size_t pitch;

    A = (float *)malloc(sizeof(float)*N*N);
    cudaMallocPitch(&dA, &pitch, sizeof(float)*N, N);

    //copy memory from unpadded array A of 760 by 760 dimensions
    //to more efficient dimensions on the device
    cudaMemcpy2D(dA,pitch,A,sizeof(float)*N,sizeof(float)*N,N,
        cudaMemcpyHostToDevice);

    ...
}
```


So..

Pitch is a good technique to speedup memory access

- There are two drawbacks that you have to live with:
 - Some wasted space
 - A bit more complicated elements access

Multi-GPU System

Nebulae: #10 in Top 500 list (June 2012)



Intel Xeon X5650 and **Nvidia GPU Tesla c2050**

Tsubame 2.0: #5 in Top 500 list



Intel Xeon X5670 and **Nvidia GPU**

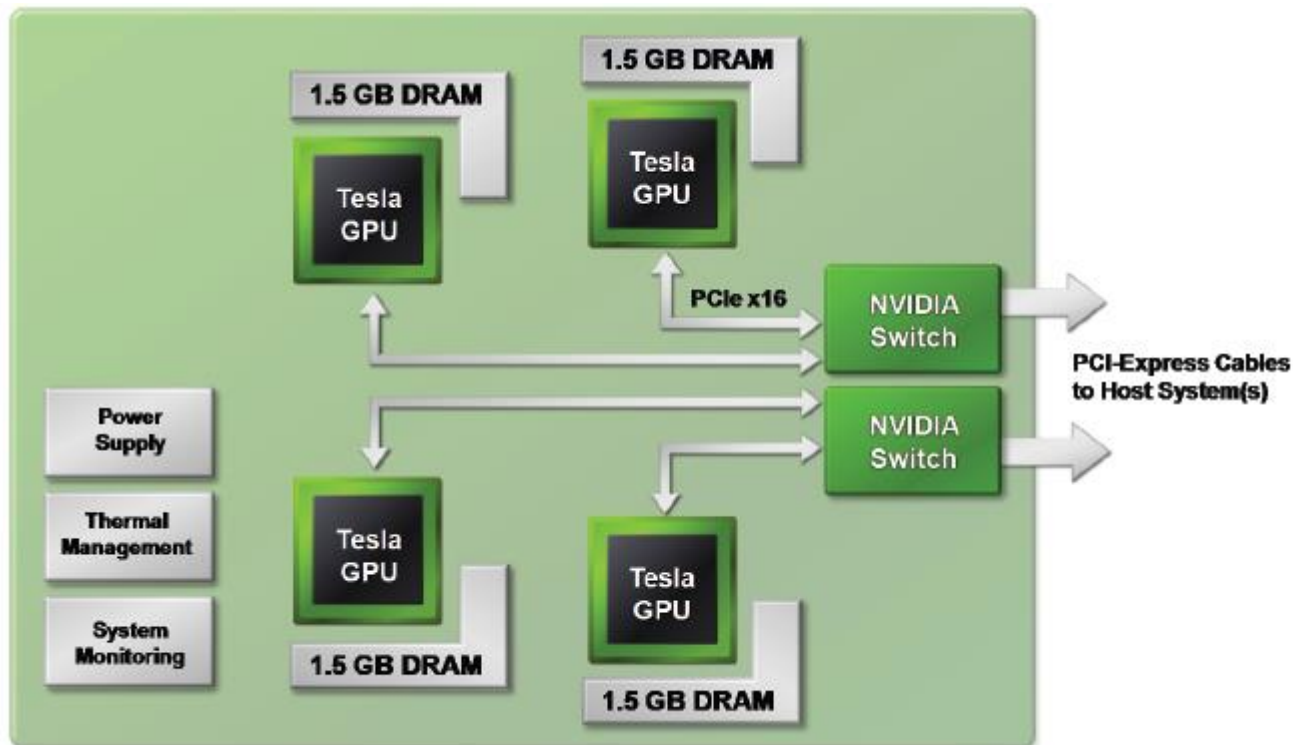
Flavors

- Multiple GPUs in the same node (e.g. PC)
- Multi-node system (e.g. MPI).



Multi-GPU configuration is here to stay!

Hardware Example: Tesla S870 Server

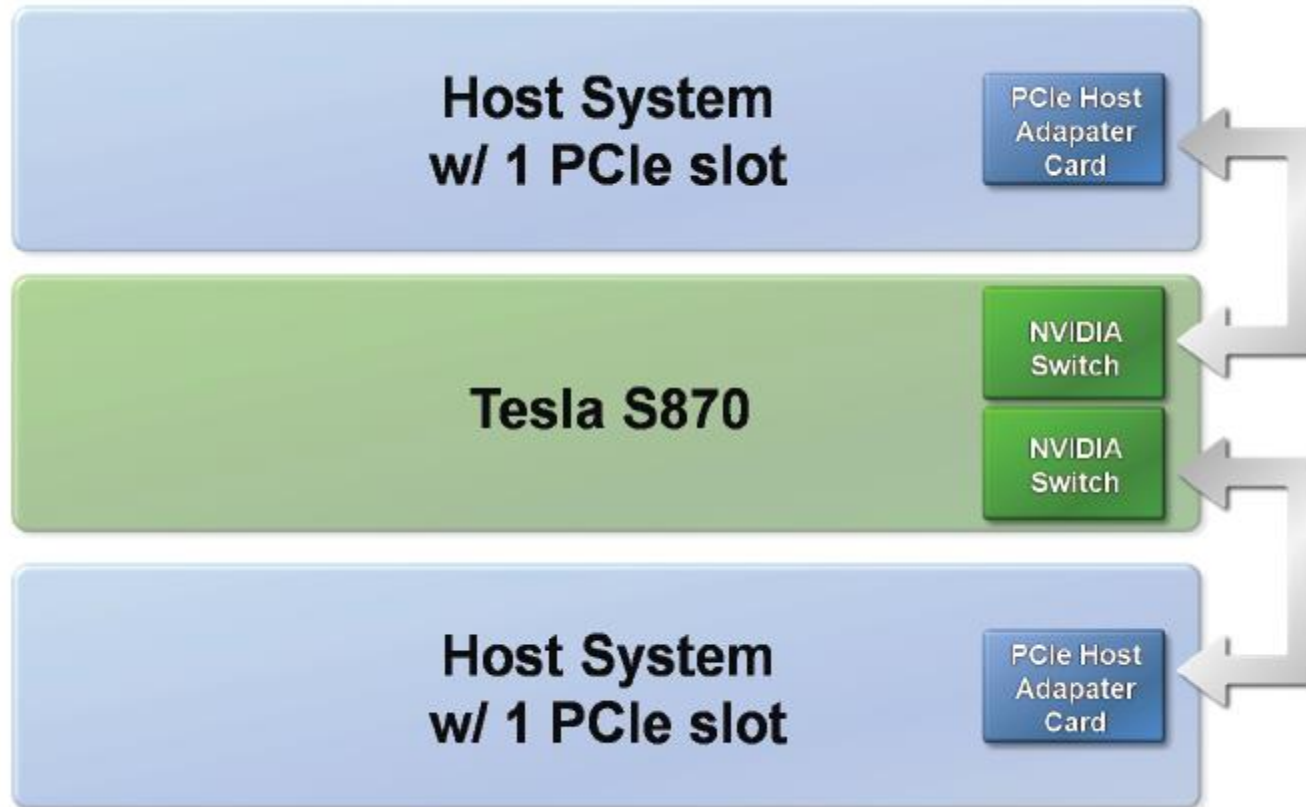


Hardware Example: Tesla S870 Server



Connected to a single-host

Hardware Example: Tesla S870 Server



Connected to a two host systems

Why Multi-GPU Solutions

- Scaling-up performance
- Another level of parallelism
- Power
- Reliability

// Run independent kernel on each CUDA device

int numDevs= 0;

cudaGetDeviceCount(&numDevs);

...

for (int d = 0; d < numDevs; d++) {

 cudaSetDevice(d);

 kernel<<<blocks, threads>>>(args);

}

CUDA Support

- `cudaGetDeviceCount(int * count)`
 - Returns in *count the number of devices
- `cudaGetDevice(int * device)`
 - Returns in *device the device on which the active host thread executes the device code.

CUDA Support

- `cudaSetDevice(devID)`
 - Device selection within the code by specifying the identifier and making CUDA kernels run on the selected GPU.

```
size_t size = 1024 * sizeof(float);
cudaSetDevice(0);           // Set device 0 as current
float* p0;
cudaMalloc(&p0, size);      // Allocate memory on device 0
MyKernel<<<1000, 128>>>(p0); // Launch kernel on device 0
cudaSetDevice(1);          // Set device 1 as current
float* p1;
cudaMalloc(&p1, size);      // Allocate memory on device 1
MyKernel<<<1000, 128>>>(p1); // Launch kernel on device 1
```

Unified Virtual Address & Peer-to-Peer Access

CUDA Support:

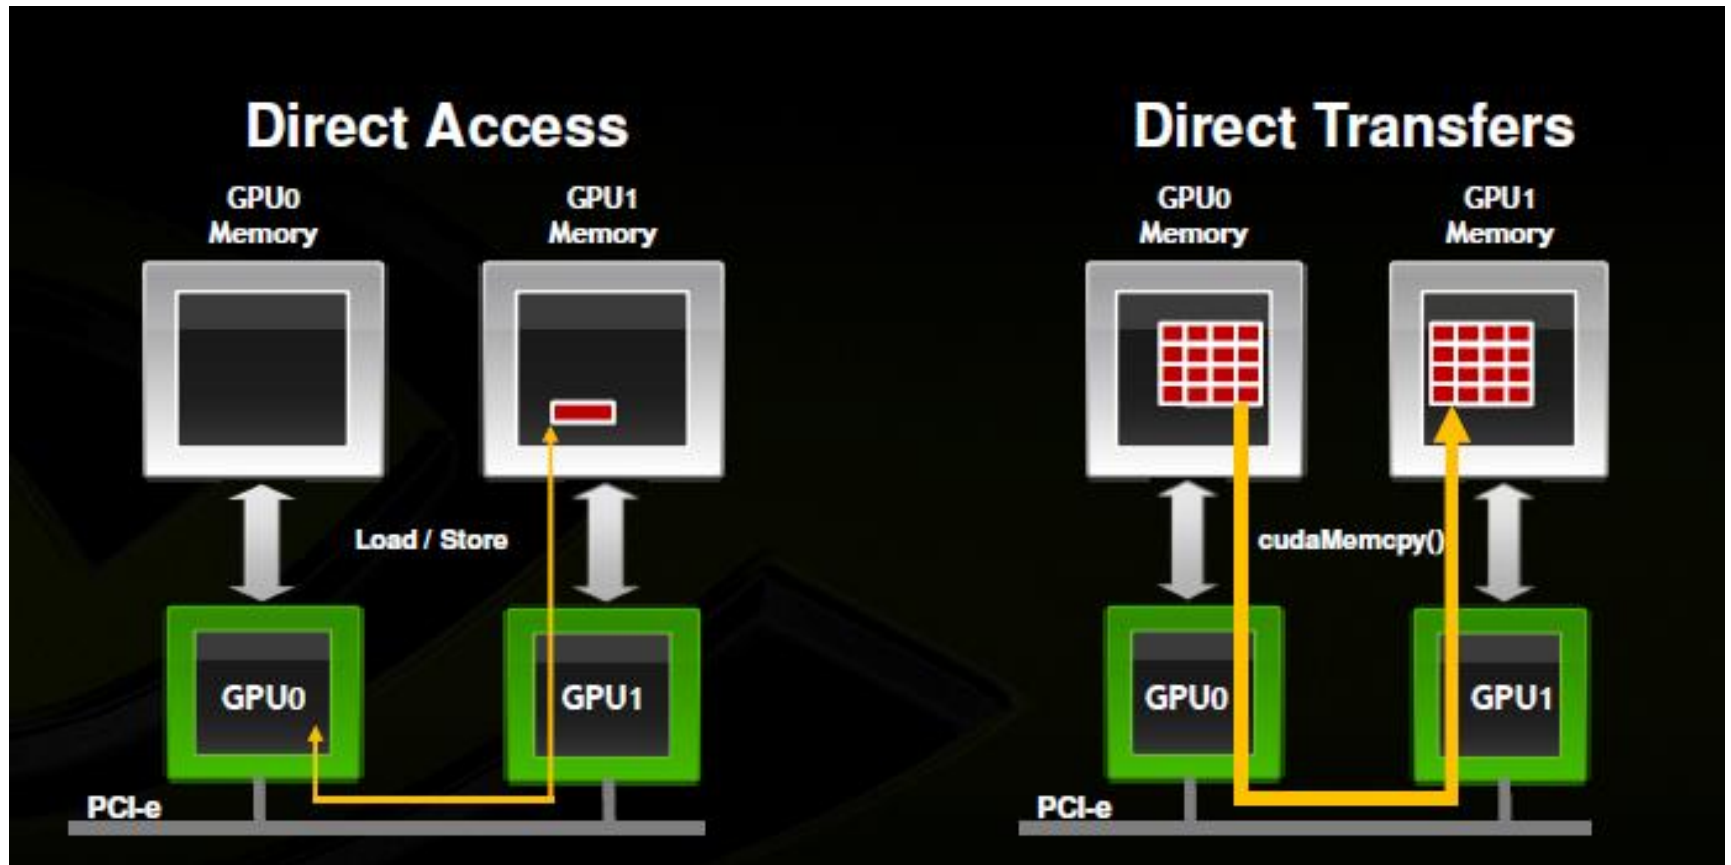
Peer to peer memory Access

- Peer-to-Peer Memory Access
 - `cudaDeviceEnablePeerAccess()` to check peer access

```
cudaSetDevice(0);           // Set device 0 as current
float* p0;
size_t size = 1024 * sizeof(float);
cudaMalloc(&p0, size);       // Allocate memory on device 0
MyKernel<<<1000, 128>>>(p0); // Launch kernel on device 0
cudaSetDevice(1);           // Set device 1 as current
cudaDeviceEnablePeerAccess(0, 0); // Enable peer-to-peer access
                                // with device 0

// Launch kernel on device 1
// This kernel launch can access memory on device 0 at address p0
MyKernel<<<1000, 128>>>(p0);
```

What we want to do ...



Does the device support P2P?

```
cudaError_t cudaDeviceCanAccessPeer  
            ( int* canAccessPeer,  
              int  device,  
              int  peerDevice)
```

- Returns 1 in canAccessPeer if device can access peerDevice.
- You need to check both directions.

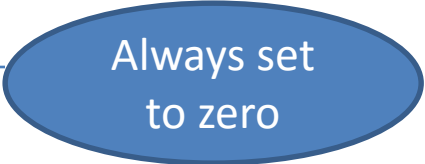
Then ...

```
cudaError_t cudaDeviceEnablePeerAccess (  
    int peerDevice,   
    unsigned int flags)
```

peerDevice
ID



Always set
to zero



Access granted by this call is **unidirectional** (i.e. current device can access peer device)

```
cudaError_t cudaDeviceDisablePeerAccess (  
    int peerDevice)
```

CUDA Support

Peer to peer memory Copy

- Using cudaMemcpyPeer()

```
cudaSetDevice(0);           // Set device 0 as current
float* p0;
size_t size = 1024 * sizeof(float);
cudaMalloc(&p0, size);      // Allocate memory on device 0
cudaSetDevice(1);          // Set device 1 as current
float* p1;
cudaMalloc(&p1, size);      // Allocate memory on device 1
cudaSetDevice(0);          // Set device 0 as current
MyKernel<<<1000, 128>>>(p0); // Launch kernel on device 0
cudaSetDevice(1);          // Set device 1 as current
cudaMemcpyPeer(p1, 1, p0, 0, size); // Copy p0 to p1
MyKernel<<<1000, 128>>>(p1); // Launch kernel on device 1
```

```
cudaMemcpyPeer ( void * dst,  
                  int dstDevice,  
                  const void * src,  
                  int srcDevice,  
                  size_t count )
```



Size of
memory
copy in
bytes

- This function is asynchronous with respect to the host.
- This function is serialized with respect to all pending and future asynchronous work into the current device.

Important: If GPU supports Unified Virtual Address,
then no need to the above function.
(We will see shortly)


Unified Virtual Address Space (UVA)

- From CUDA 4.0
- puts all CUDA execution, CPU and GPU, in the **same address space**
- Requires Fermi-class GPU and above
- Requires 64-bit application
- Call **cudaGetDeviceProperties()** for all participating devices and check **unifiedAddressing** flag


Unified Virtual Addressing

Easier to Program with Single Address Space

No UVA: Multiple Memory Spaces



UVA : Single Address Space



Easier Memory Copy

- **Between host and multiple devices:**

`cudaMemcpy(gpu0_buf, host_buf, buf_size, cudaMemcpyDefault)`

`cudaMemcpy(gpu1_buf, host_buf, buf_size, cudaMemcpyDefault)`

`cudaMemcpy(host_buf, gpu0_buf, buf_size, cudaMemcpyDefault)`

`cudaMemcpy(host_buf, gpu1_buf, buf_size, cudaMemcpyDefault)`

- **Between two devices:**

`cudaMemcpy(gpu0_buf, gpu1_buf, buf_size, cudaMemcpyDefault)`

- `cudaMemcpy()` knows that our buffers are on different devices
- (UVA), will do a P2P copy
- Note that this will transparently fall back to a normal copy through the host if P2P is not available

Example:

Direct N-Body

- Simulation of dynamical system of N-bodies
- $O(N^2)$
- Compute-Bound application
- Assume we have K GPUs
 - So each GPU is responsible for N/K bodies
- For each iteration:
 - Get all N up-to-date positions onto each GPU
 - Compute accelerations (N/k per GPU)
 - Integrate position, velocity (N/k per GPU)

Example:

Direct N-Body

- Sharing data among GPUs: options
 - **Explicit copies via host**
 - Zero-copy shared host array
(`cudaMallocHost()`)
 - Per-device arrays with peer-to-peer exchange transfers (UVA)
 - Peer-to-peer memory access

N-Body

Explicit Copy Via Host

```
for(;;) {  
    for (int d = 0; d < devs; d++) {  
        cudaSetDevice(d);  
        cudaMemcpyAsync(pos[d], in, devs*bytes, H2D, s[d]);  
    }  
    for (int d = 0; d < devs; d++) {  
        cudaSetDevice(d);  
        integrate<<<b, t, 0, s[d]>>>(pos[d], otherArgs);  
    }  
    for (int d = 0; d < devs; d++) {  
        cudaSetDevice(d);  
        cudaMemcpyAsync(&out[offset[d]], pos[d], bytes, D2H,  
            s[d]);  
    }  
}
```

Example:

Direct N-Body

- Sharing data among GPUs: options
 - Explicit copies via host
 - **Zero-copy shared host array** (direct device access to host memory, through PCIe, which is slow) ... `cudaMallocHost()` or `cudaHostAlloc()` ... so, use it when:
 - You copy data to the device and access it there only once AND/OR
 - You generate data on the device and copy back to host without reuse AND/OR
 - Your kernel(s) that access the memory are compute bound
 - Per-device peer-to-peer exchange transfers (UVA)
 - Peer-to-peer memory access

N-Body

Zero-copy

```
// Create input and output arrays  
cudaHostAlloc(&in, bytes, cudaHostAllocMapped |  
cudaHostAllocPortable);  
cudaHostAlloc(&out, bytes, cudaHostAllocMapped  
| cudaHostAllocPortable);
```

Allocates size bytes of host memory that is **page-locked** and accessible to the device.

**Important: If GPU supports Unified Virtual Address,
then no need to the above function.
(We will see shortly)**

N-Body

Zero-copy

```
// Create input and output arrays
cudaHostAlloc(&in, bytes, cudaHostAllocMapped |
cudaHostAllocPortable);
cudaHostAlloc(&out, bytes, cudaHostAllocMapped
| cudaHostAllocPortable);

for (int d = 0; d < devCount; d++) {
    cudaSetDevice(d);
    cudaHostGetDevicePointer(&dout[d], hostPtr, 0);
    cudaHostGetDevicePointer(&din[d], hostPtr, 0);
}
```



pointer that will be passed to the device to access host memory

Example:

Direct N-Body

- Sharing data among GPUs: options
 - Explicit copies via host
 - Zero-copy shared host array
(`cudaMallocHost()`)
 - **Per-device peer-to-peer exchange transfers**
 - UVA as we have seen
 - Non-UVA:
 - `cudaMemcpyPeer()`
 - Copies memory from one device to memory on another device
 - Peer-to-peer memory access

Example:


Direct N-Body

- Sharing data among GPUs: options
 - Explicit copies via host
 - Zero-copy shared host array
(`cudaMallocHost()`)
 - Per-device peer-to-peer exchange transfers
 - **Peer-to-peer memory access**
 - Pass pointer to memory on device A to kernel running on device B
 - Requires UVA
 - Must first enable peer access for every pair:
 - `cudaDeviceEnablePeerAccess`

Dynamic Parallelism


The Usual case

- Data travels back and forth between the CPU and GPU many times.
- Reason: because of the inability of the GPU to create more work on itself depending on the data.



With Dynamic Parallelism:

- GPU can generate work on itself **without involvement of CPU.**
- Permits Dynamic Runtime decisions.
- Kernels can start new kernels
- Streams can spawn new streams.



CUDA 5.0 and later on devices of Compute Capability 3.5 or higher

```

__global__ ChildKernel(void* data){
    //Operate on data
}

__global__ ParentKernel(void *data){
    if (threadIdx.x == 0) {
        ChildKernel<<<1, 32>>>(data);
        cudaThreadSynchronize();
    }
    syncthreads();
    //Operate on data
}

// In Host Code
ParentKernel<<<8, 32>>>(data);

```


A kernel can call another kernel that call another kernel up to 24 nested ...
 Subject to the availability of resources.

Important

- As on the host, device kernel launch is asynchronous.
- Successful execution of a kernel launch means that the kernel is queued;
 - it may begin executing immediately,
 - or it may execute later when resources become available.
- Note that every thread that encounters a kernel launch executes it. So be careful!
- Child grids always complete before the parent grids that launch them, even if there is no explicit synchronization.

Important

- The CUDA Device Runtime guarantees that parent and child grids have **a fully consistent view of global memory when the child starts and ends.**



Important

- By default, grids launched within a thread block are executed sequentially.
- This happens even if grids are launched by different threads within the block.
- To deal with this drawback → streams
- streams created on the host cannot be used on the device.
- Streams created in a block can be used by all threads in that block.

```
cudaStream_t s;
```

```
cudaStreamCreateWithFlags(&s, cudaStreamNonBlocking);
```

Important

- If the parent kernel needs results computed by the child kernel to do its own work → it must ensure that the child grid has finished execution before continuing
 - by explicitly synchronizing using `cudaDeviceSynchronize(void)`.
 - This function waits for completion of all grids previously launched by the thread block from which it has been called.

Example


```
void threadBlockDeviceSynchronize(void)
{
    __syncthreads();
    if(threadIdx.x == 0)
        cudaDeviceSynchronize();
    __syncthreads();
}
```

—————→ To ensure all launches
have been made.

What do we gain?

- Reduction in trips to CPU
- Recursion
- More freedom where data generated by the kernel decides how to partition the data for lower-level of the hierarchy.

How to Compile and Link?



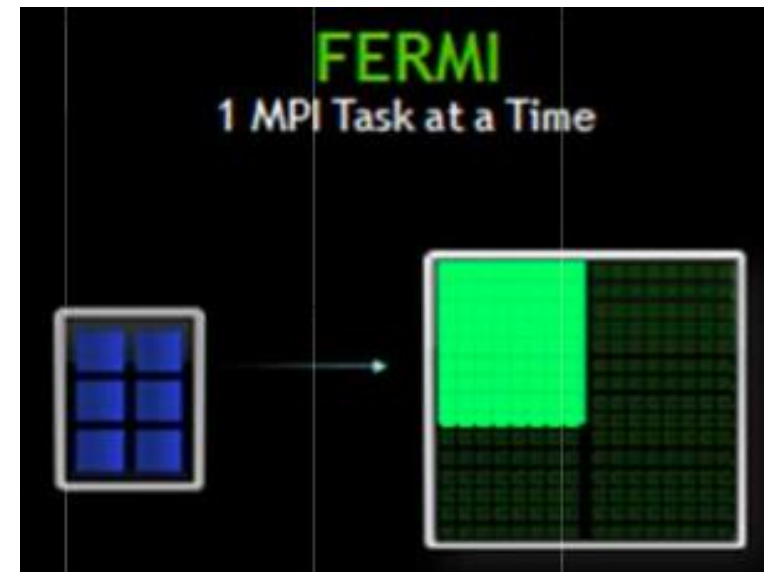
`nvcc -arch=sm_35 -rdc=true myprog.cu -lcudadevrt`

generate relocatable device code, required for later linking

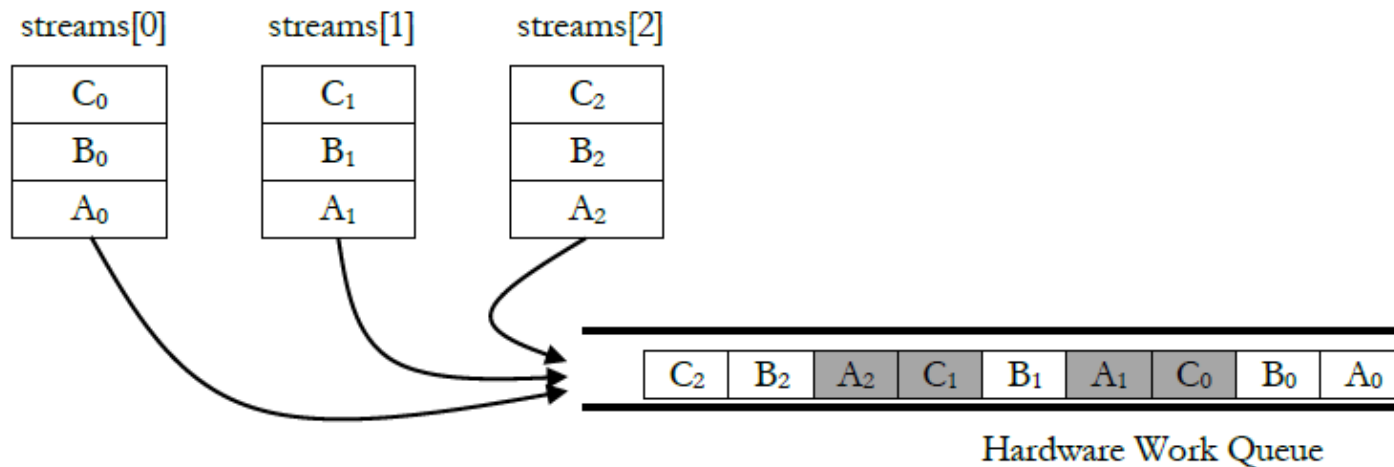

Hyper-Q

Till Fermi

- Only one work queue
- Even though Fermi allows 16 concurrent kernels.
- GPU resources not fully utilized



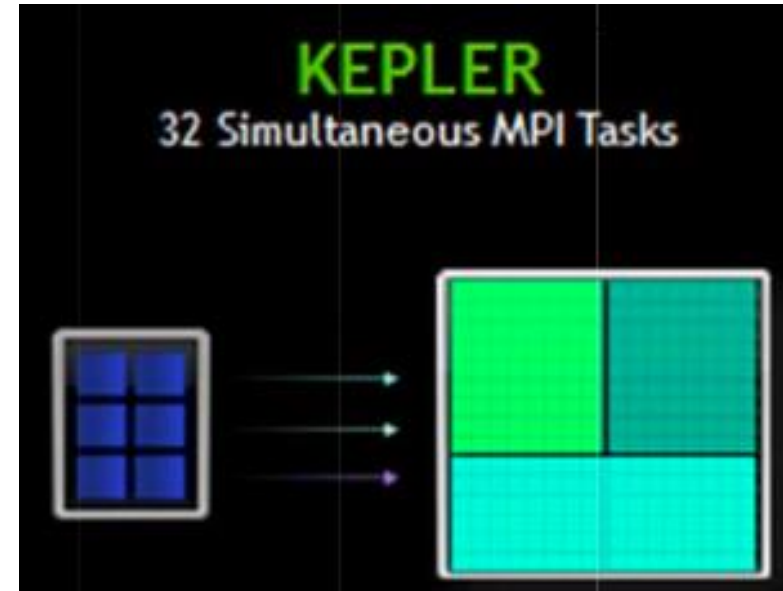
Fermi already supported 16 way concurrency of kernel launches from separate streams
Pending work is bottlenecked on 1 work queue.
GPU's computational resources not being utilized fully.

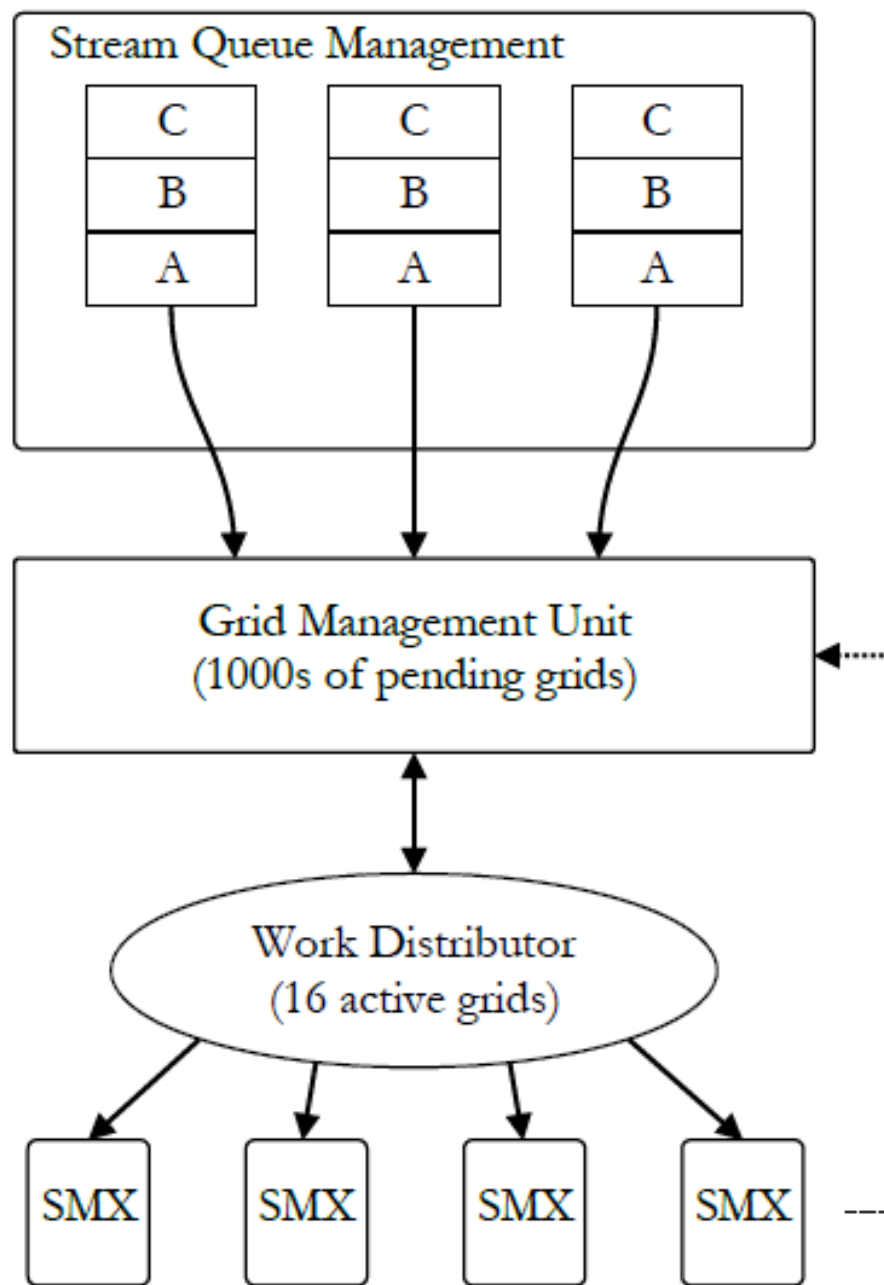


With Hyper-Q

- Starting with Kepler
- We can have connection from multiple CUDA streams, Message Passing Interface (MPI) processes, or multiple threads of the same process.
 - 32 concurrent work queues, can receive work from 32 process cores at the same time.
 - 3X Performance increase on Fermi

With Hyper-Q





Conclusions

- There are many performance enhancement techniques in our arsenal:
 - Alignment
 - Streams
 - Pinned pages
 - Asynchronous execution
 - Dynamic Parallelism
 - Multi-GPU
- If your program is making use of a lot of FP operations, be careful about rounding errors.
- There are tools to help you!