



CSCI-GA.3033-004

Graphics Processing Units (GPUs): Architecture and Programming

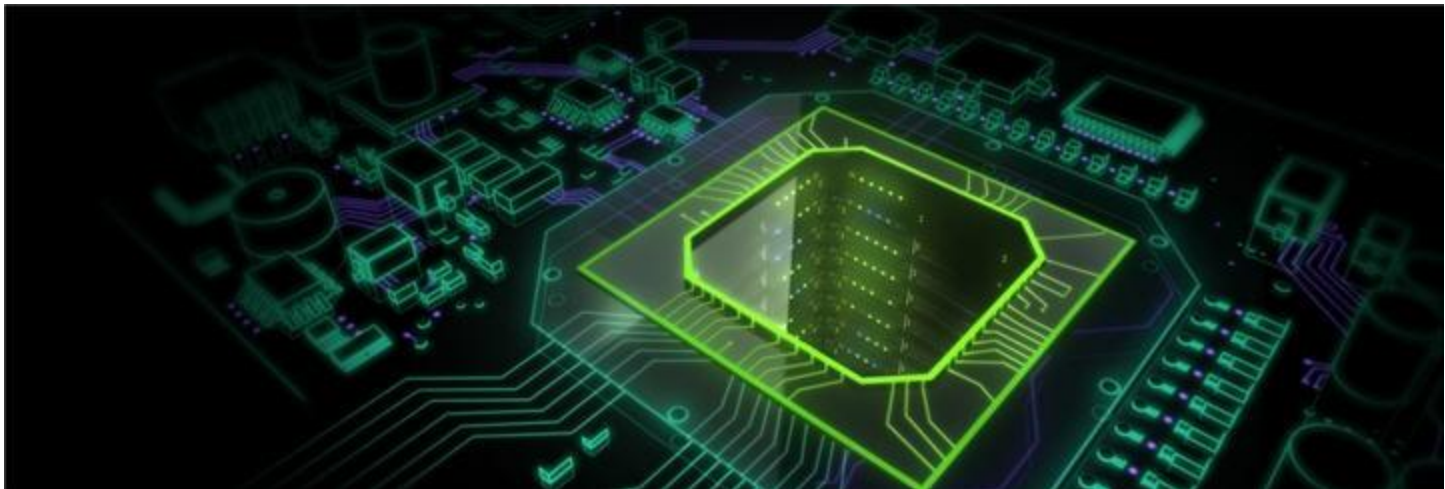
CUDA

Advanced Techniques 3

Mohamed Zahran (aka Z)

mzahran@cs.nyu.edu

<http://www.mzahran.com>



In This Lecture ...

- More about performance
- Parallel Patterns
- Error Handling

More About Performance

Hardware configuration can be safely ignored when designing a software for correctness but must be considered in the code structure when designing for peak performance.

Throughput


Latency

Some Insights About Performance

Occupancy





Utilization

It is a common belief that ...





- More threads is better
 - because it needs more threads hide latency
- 
- ```
graph TD; A(latency) --> B(Computational); A --> C(memory);
```
- The diagram shows the word "latency" from the previous list item circled in red. Two red arrows originate from the bottom of the circle, pointing downwards and outwards to the words "Computational" on the left and "memory" on the right.

But is it always true?

Multiplication of two large matrices, single precision (SGEMM):

|                   | CUBLAS 1.1                                                                            | CUBLAS 2.0                                                                              |                          |
|-------------------|---------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------|--------------------------|
| Threads per block | 512  | 64   | 8x smaller thread blocks |
| Occupancy (G80)   | 67%  | 33%  | 2x lower occupancy       |
| Performance (G80) | 128 Gflop/s                                                                           | 204 Gflop/s                                                                             | 1.6x higher performance  |

Batch of 1024-point complex-to-complex FFTs, single precision:

|                   | CUFFT 2.2                                                                              | CUFFT 2.3                                                                                |                          |
|-------------------|----------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------|--------------------------|
| Threads per block | 256  | 64   | 4x smaller thread blocks |
| Occupancy (G80)   | 33%  | 17%  | 2x lower occupancy       |
| Performance (G80) | 45 Gflop/s                                                                             | 93 Gflop/s                                                                               | 2x higher performance    |

# Latency Vs Throughput


- Latency (how much time) is **time**
  - instruction takes 4 cycles per warp
  - memory takes 400 cycles
- Throughput (how many operations per cycle or second) is **rate**
  - Arithmetic:  $1.3 \text{ Tflop/s} = 480 \text{ ops/cycle}$  (op=multiply-add)
  - Memory:  $177 \text{ GB/s} \approx 32 \text{ ops/cycle}$  (op=32-bit load)

# Hide Latency is ...

- Doing other operations while waiting
- This will make the kernel runs faster
- But not at the peak performance
- What can we do?



# Little's Law



Needed parallelism = **Latency** x **Throughput**

# Examples from GPU

| GPU model | Latency<br>(cycles) | Throughput<br>(cores/SM) | Parallelism<br>(operations/SM) |
|-----------|---------------------|--------------------------|--------------------------------|
| G80-GT200 | $\approx 24$        | 8                        | $\approx 192$                  |
| GF100     | $\approx 18$        | 32                       | $\approx 576$                  |
| GF104     | $\approx 18$        | 48                       | $\approx 864$                  |

Average latency of a computational operation



Less operations means idle cycle



# So ...

- Higher performance does not mean more threads but **higher utilization**
- Utilization is related to **parallelism**
- We can increase utilization by
  - increasing throughput
    - Instruction level parallelism
    - Thread level parallelism
  - decreasing latency

**Occupancy is not utilization, but one of the contributing factors.**

# Occupancy Calculator API

```
__global__ void MyKernel(int *d, int *a, int *b) {
 int idx = threadIdx.x + blockIdx.x * blockDim.x;
 d[idx] = a[idx] * b[idx]; }
```

```
int main() {
 int numBlocks;
 int blockSize = 32;
 int device;
 cudaDeviceProp prop;
 int activeWarps;
 int maxWarps;

 cudaGetDevice(&device);
 cudaGetDeviceProperties(&prop, device);

 cudaOccupancyMaxActiveBlocksPerMultiprocessor(
 &numBlocks,
 MyKernel,
 blockSize,
 0);

 activeWarps = numBlocks * blockSize / prop.warpSize;
 maxWarps = prop.maxThreadsPerMultiProcessor / prop.warpSize;
}
```

## cudaOccupancyMaxActiveBlocksPerMultiprocessor

- From CUDA 6.5
- Produces an occupancy prediction based on:
  - the block size
  - shared memory usage of a kernel
- Reports occupancy in terms of the number of concurrent thread blocks per multiprocessor
- Don't forget: it is just a prediction!
- Arguments:
  1. pointer to an integer (where #blocks will be reported)
  2. kernel
  3. block size
  4. dynamic shared memory per block in bytes

# How about memory?

maximizing overall memory throughput for  
the application

=

minimize data transfers  
with low bandwidth

host  $\leftrightarrow$  device

Global mem access

# This means ...Typically

1. Load data from device memory to shared memory.
2. Synchronize with all the other threads of the block so that each thread can safely read shared memory locations that were populated by different threads.
3. Process the data in shared memory.
4. Synchronize again if necessary to make sure that shared memory has been updated with the results.
5. Write the results back to device memory.

But accessing global memory is a  
necessary evil ... So:

- Can we apply the same technique (i.e. Little's law) to memory?

Needed parallelism = **Latency** x **Throughput**

|            | <b>Latency</b>   | <b>Throughput</b> | <b>Parallelism</b> |
|------------|------------------|-------------------|--------------------|
| Arithmetic | ≈18 cycles       | 32 ops/SM/cycle   | 576 ops/SM         |
| Memory     | < 800 cycles (?) | < 177 GB/s        | < 100 KB           |

This means that to hide memory latency you need to keep 100KB in flight.  
But less if the kernel is compute bound!



# How Can You Get 100KB From Threads?

- Use more threads
- Use more instructions per thread
- Use more data per thread

# Now for some commonly used parallel patterns

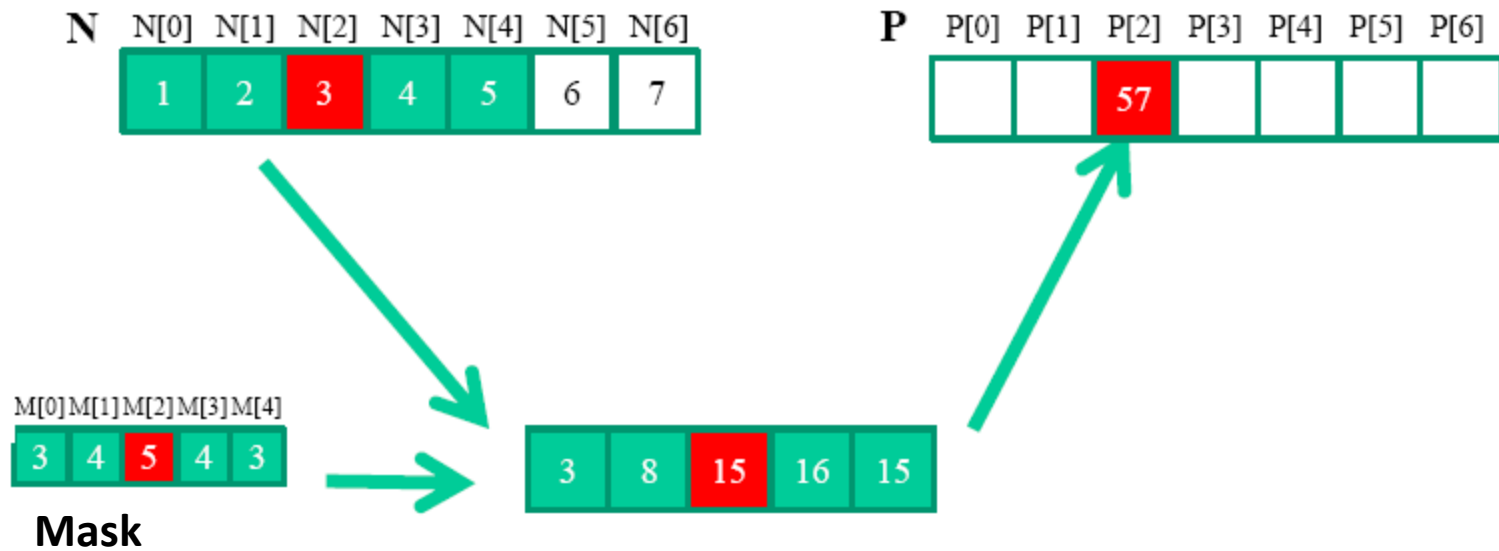
- Convolution
- Reduction tree
- Prefix sum

Pattern: Convolution

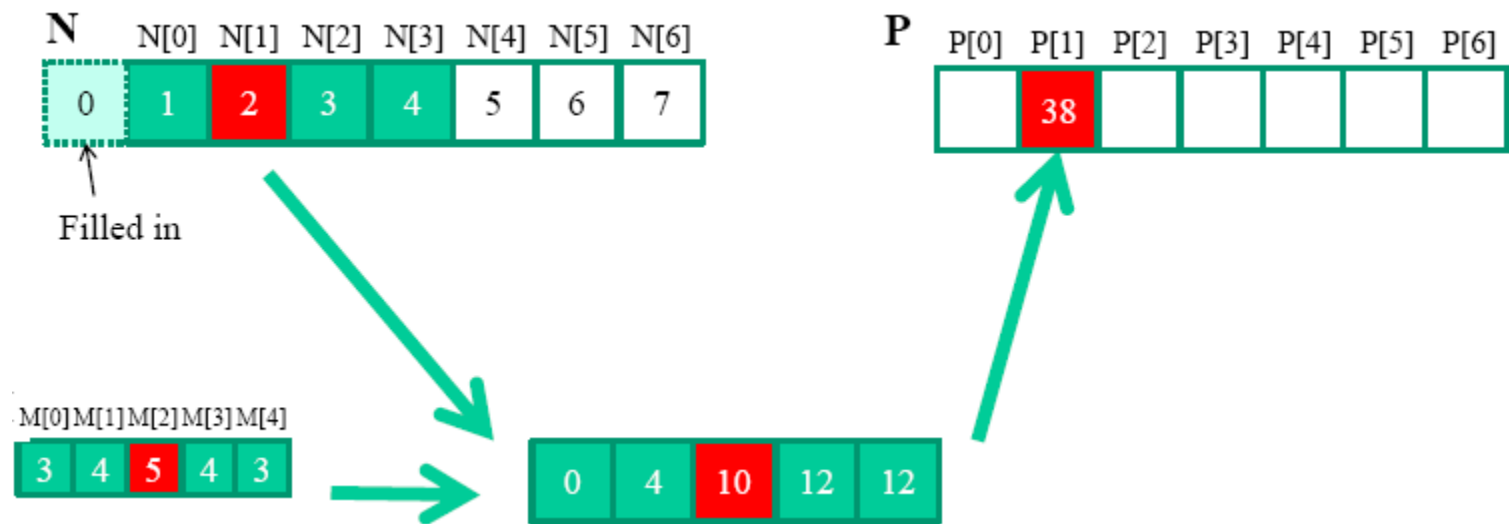
# Convolution

- An Array operation
- Output data element = weighted sum of a collection of neighboring input elements.
- The weights are defined by an input mask array.
- Usually used as filters to transform signals (or pixels or ...) into more desirable form.

# Convolution



# Convolution



Convolution can also be 2D.

# Convolution

N


|   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|
| 1 | 2 | 3 | 4 | 5 | 6 | 7 |
| 2 | 3 | 4 | 5 | 6 | 7 | 8 |
| 3 | 4 | 5 | 6 | 7 | 8 | 9 |
| 4 | 5 | 6 | 7 | 8 | 5 | 6 |
| 5 | 6 | 7 | 8 | 5 | 6 | 7 |
| 6 | 7 | 8 | 9 | 0 | 1 | 2 |
| 7 | 8 | 9 | 0 | 1 | 2 | 3 |

P

|  |  |     |  |  |  |  |
|--|--|-----|--|--|--|--|
|  |  |     |  |  |  |  |
|  |  |     |  |  |  |  |
|  |  | 321 |  |  |  |  |
|  |  |     |  |  |  |  |
|  |  |     |  |  |  |  |
|  |  |     |  |  |  |  |
|  |  |     |  |  |  |  |

M

|   |   |   |   |   |
|---|---|---|---|---|
| 1 | 2 | 3 | 2 | 1 |
| 2 | 3 | 4 | 3 | 2 |
| 3 | 4 | 5 | 4 | 3 |
| 2 | 3 | 4 | 3 | 2 |
| 1 | 2 | 3 | 2 | 1 |



|   |    |    |    |    |
|---|----|----|----|----|
| 1 | 4  | 9  | 8  | 5  |
| 4 | 9  | 16 | 15 | 12 |
| 9 | 16 | 25 | 24 | 21 |
| 8 | 15 | 24 | 21 | 16 |
| 5 | 12 | 21 | 16 | 5  |

# Convolution

```
__global__ void convolution_1D_basic_kernel(float *N, float *M, float *P,
int Mask_Width, int Width) {

 int i = blockIdx.x*blockDim.x + threadIdx.x;

 float Pvalue = 0;
 int N_start_point = i - (Mask_Width/2);
 for (int j = 0; j < Mask_Width; j++) {
 if (N_start_point + j >= 0 && N_start_point + j < Width) {
 Pvalue += N[N_start_point + j]*M[j];
 }
 }
 P[i] = Pvalue;
}
```

The 1D Version

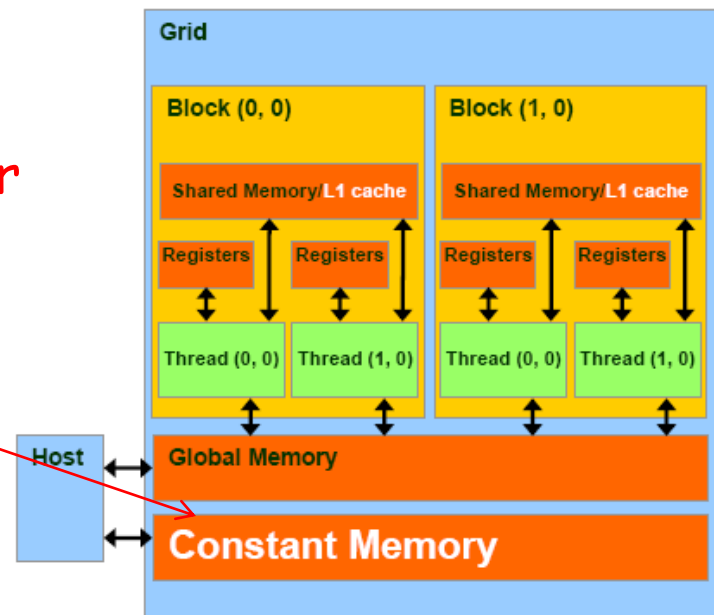
- Thread organized as 1D grid.
- Pvalue allows intermediate values to be accumulated in registers to save DRAM bw.
- We assume *ghost* values are 0.
- There will be control flow divergence (due to ghost elements).
- **Ratio of floating point arithmetic calculation to global memory access is  $\sim 1.0 \rightarrow$  What can we do??**



# Regarding Mask M

- Size of  $M$  is typically small.
- The contents of  $M$  do not change during execution.
- All threads need to access  $M$  and in the same order.

Doesn't this make  $M$  a good candidate for constant memory?



# Constant Memory

- Constant memory variables are visible to all thread blocks.
- Constant memory variables cannot be changed during kernel execution.
- The size of constant memory can vary from device to device.

# Mask M and Constant Memory

- In host:
  - `#define MASK_WIDTH 10`
  - `__constant__ float M[MASK_WIDTH]`
  - Allocate and initialize a mask `h_M`
  - `cudaMemcpyToSymbol(M, h_M, MASK_WIDTH * sizeof(float),  
offset, kind);`
- Kernel functions
  - access constant memory variables as global variables → no need to pass pointers of these variables to the kernel as parameter.

Question: Isn't the constant memory also in DRAM? Why is it assumed faster than global memory?

Answer:

- CUDA runtime knows that constant memory variables are not modified.
- It directs the hardware to aggressively **cache** them during kernel execution.

Pattern: Reduction Tree

# What is it?

- A commonly used strategy for processing large input data sets
  - There is no required order of processing elements in a data set
  - Partition the data set into smaller chunks
  - Have each thread to process a chunk
  - Use a reduction tree to summarize the results from each chunk into the final answer
- Google and Hadoop MapReduce frameworks are examples of this pattern

# What is it?


- Summarize a set of input values into one value using a “reduction operation”
  - Max
  - Min
  - Sum
  - Product
- Often with user defined reduction operation function as long as the operation
  - Is associative and commutative
  - Has a well-defined identity value (e.g., 0 for sum)

## An efficient sequential reduction algorithm performs $N$ operations in $O(N)$

- Initialize the result as an identity value for the reduction operation
  - Smallest possible value for max reduction
  - Largest possible value for min reduction
  - 0 for sum reduction
  - 1 for product reduction
- Scan through the input and perform the reduction operation between the result value and the current input value



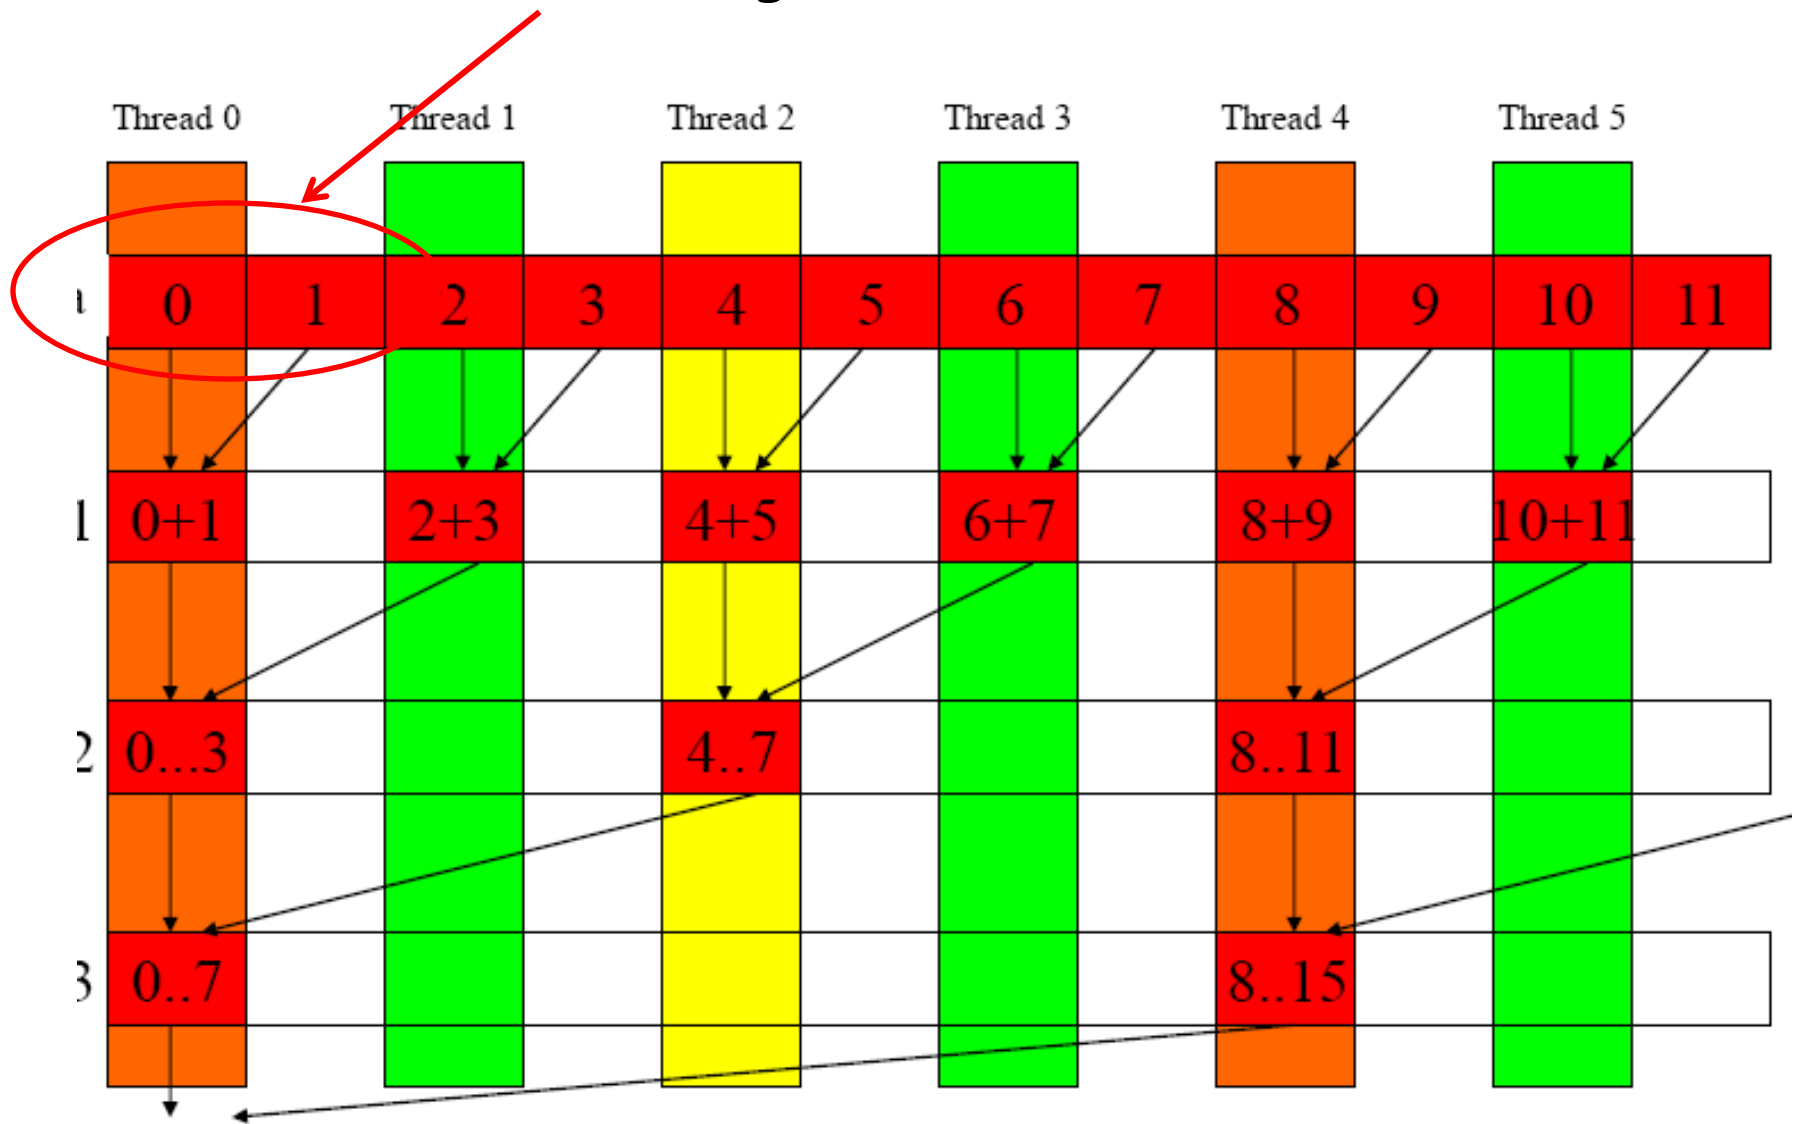
A parallel reduction tree algorithm performs  $N-1$  Operations in  $\log(N)$  steps



# Straightforward Implementation


- The original vector is in device global memory
- The shared memory is used to hold a partial sum vector
- Each step brings the partial sum vector closer to the sum
- The final sum will be in element 0
- Reduces global memory traffic due to partial sum values

A lot of branch divergence.



Thread 0 Thread 1 Thread 2

Thread 14 Thread 15



A Better Version

# Be Careful!

- Although the number of "operations" is  $N$ , each "operation involves much more complex address calculation and intermediate result manipulation.
- If the parallel code is executed on a single-thread hardware, it would be significantly slower than the code based on the original sequential algorithm.

Pattern: Prefix Sum (Scan)

# Scan / Parallel Prefix Sum

|   |   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|---|
| 3 | 1 | 7 | 0 | 4 | 1 | 6 | 3 |
|---|---|---|---|---|---|---|---|



|   |   |   |    |    |    |    |    |
|---|---|---|----|----|----|----|----|
| 0 | 3 | 4 | 11 | 11 | 15 | 16 | 22 |
|---|---|---|----|----|----|----|----|

- Given an array  $A = [a_0, a_1, \dots, a_{n-1}]$   
and a binary associative operator  $@$  with identity  $I$ 
  - $\text{scan}(A) = [I, a_0, (a_0 @ a_1), \dots, (a_0 @ a_1 @ \dots @ a_{n-2})]$
- This is the **exclusive scan**

# Inclusive Scan

|   |   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|---|
| 3 | 1 | 7 | 0 | 4 | 1 | 6 | 3 |
|---|---|---|---|---|---|---|---|



|   |   |    |    |    |    |    |    |
|---|---|----|----|----|----|----|----|
| 3 | 4 | 11 | 11 | 15 | 16 | 22 | 25 |
|---|---|----|----|----|----|----|----|

- Given an array  $A = [a_0, a_1, \dots, a_{n-1}]$  and a binary associative operator  $@$  with identity  $I$ 
  - $\text{scan}(A) = [a_0, (a_0 @ a_1), \dots, (a_0 @ a_1 @ \dots @ a_{n-1})]$



# Why?

- Scan is used as a building block for many parallel algorithms
  - Radix sort
  - Quicksort
  - String comparison
  - Lexical analysis
  - Run-length encoding
  - Histograms
  - Etc.


# A Inclusive Scan Application Example

- Assume that we have a 100-inch sausage to feed 10
- We know how much each person wants in inches
  - [3 5 2 7 28 4 3 0 8 1]
- How do we cut the sausage quickly?
- How much will be left
- Method 1: cut the sections sequentially: 3 inches first, 5 inches second, 2 inches third, etc.
- Method 2: calculate Prefix scan
  - [3, 8, 10, 17, 45, 49, 52, 52, 60, 61] (39 inches left)

# Other Examples

- Assigning camp slots
- Assigning farmer market space
- Allocating memory to parallel threads
- Allocating memory buffer for communication channels
- ...

# Sequential algorithm



```
void scan(float* output, float* input, int length)
{
 output[0] = 0;
 for(int j = 1; j < length; ++j)
 {
 output[j] = input[j-1] + output[j-1];
 }
}
```

- $N$  additions
- Use a guide:
  - Want parallel to be *work efficient*
  - Does similar amount of work


# A Parallel Inclusive Scan Algorithm

|    |   |   |   |   |   |   |   |   |
|----|---|---|---|---|---|---|---|---|
| T0 | 3 | 1 | 7 | 0 | 4 | 1 | 6 | 3 |
|----|---|---|---|---|---|---|---|---|

1. Read input from device memory to shared memory

Each thread reads one value from the input array in device memory into shared memory array.

# A Parallel Scan Algorithm




1. (previous slide)
2. Iterate  $\log(n)$  times: Threads *stride* to  $n$ : Add pairs of elements *stride* elements apart. Double *stride* at each iteration. (note must double buffer shared mem arrays)

Iteration #1  
Stride = 1

- Active threads: *stride* to  $n-1$  ( $n$ -*stride* threads)
- Thread  $j$  adds elements  $j$  and  $j$ -*stride* from T0 and writes result into shared memory buffer T1 (ping-pong)


# A Parallel Scan Algorithm



1. Read input from device memory to shared memory.
2. Iterate  $\log(n)$  times: Threads *stride* to *n*: Add pairs of elements *stride* elements apart. Double *stride* at each iteration. (note must double buffer shared mem arrays)

Iteration #2  
Stride = 2

# A Parallel Scan Algorithm



Iteration #3  
Stride = 4

1. Read input from device memory to shared memory. Set first element to zero and shift others right by one.
2. Iterate  $\log(n)$  times: Threads *stride* to *n*: Add pairs of elements *stride* elements apart. Double *stride* at each iteration. (note must double buffer shared memory arrays)
3. Write output from shared memory to device memory




# Work Efficiency Considerations

- The first-attempt Scan executes  $\log(n)$  iterations
- This scan algorithm is not very work efficient
  - Sequential scan algorithm does  $n$  adds
  - A factor of  $\log(n)$  hurts: 20x for  $10^6$  elements!
- A parallel algorithm can be slow when execution resources are saturated due to low work efficiency

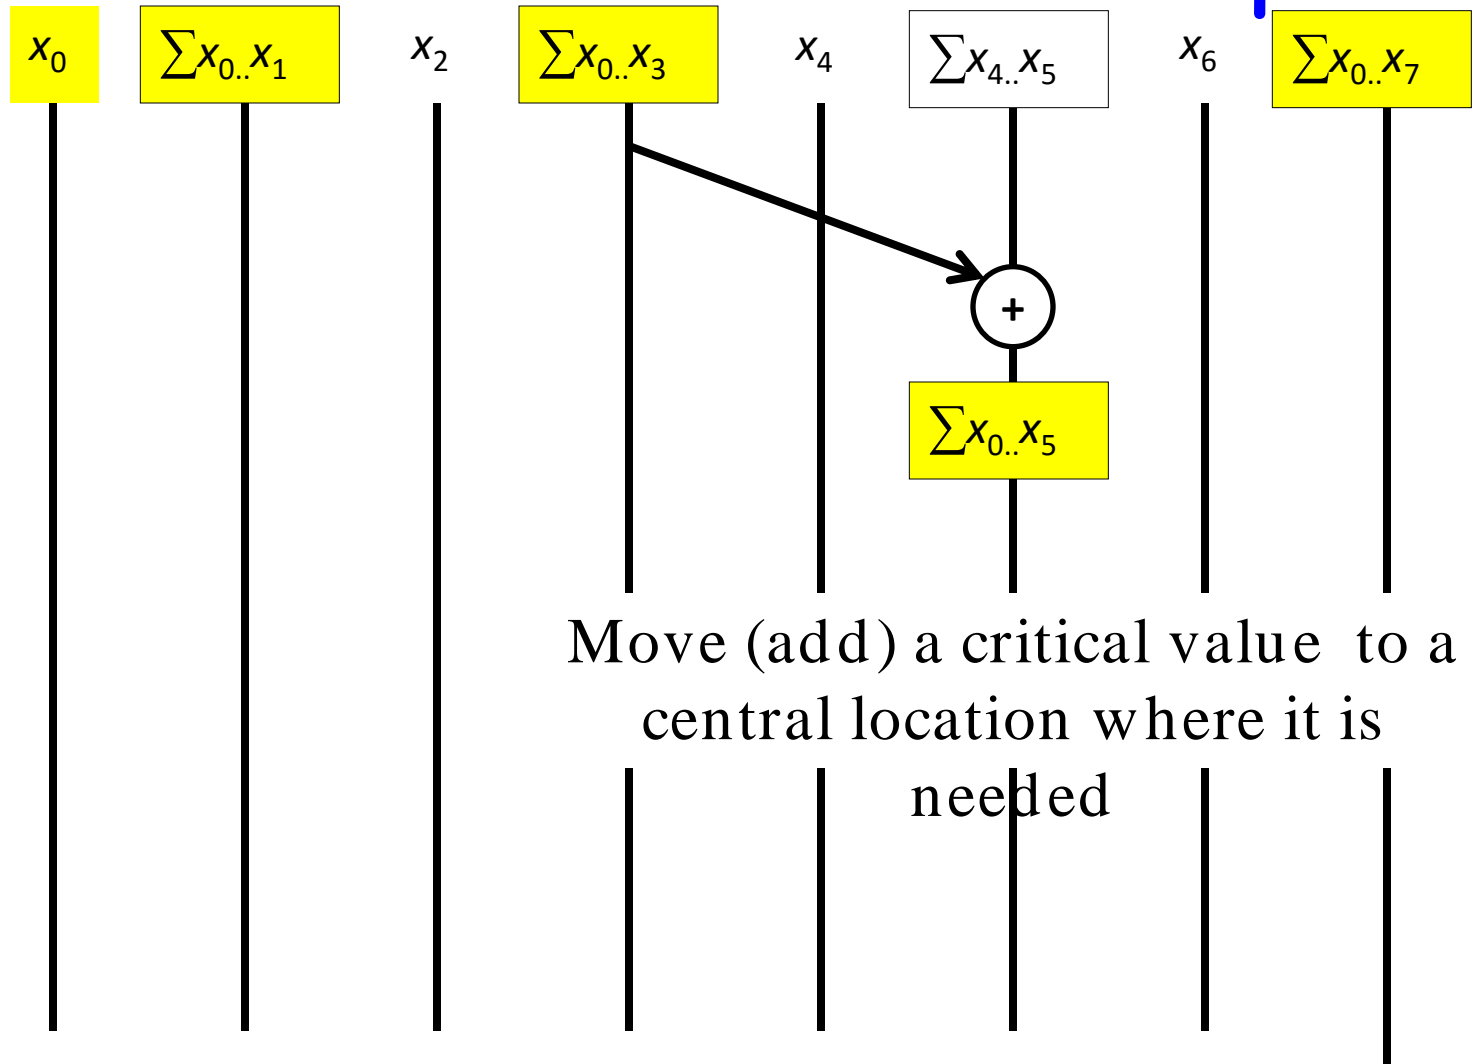
# Improving Efficiency

- A common parallel algorithm pattern:
  - Balanced Trees*
  - Build a balanced binary tree on the input data and sweep it to and from the root
  - Tree is not an actual data structure, but a concept to determine what each thread does at each step
- For scan:
  - Traverse down from leaves to root building partial sums at internal nodes in the tree
    - Root holds sum of all leaves
  - Traverse back up the tree building the scan from the partial sums

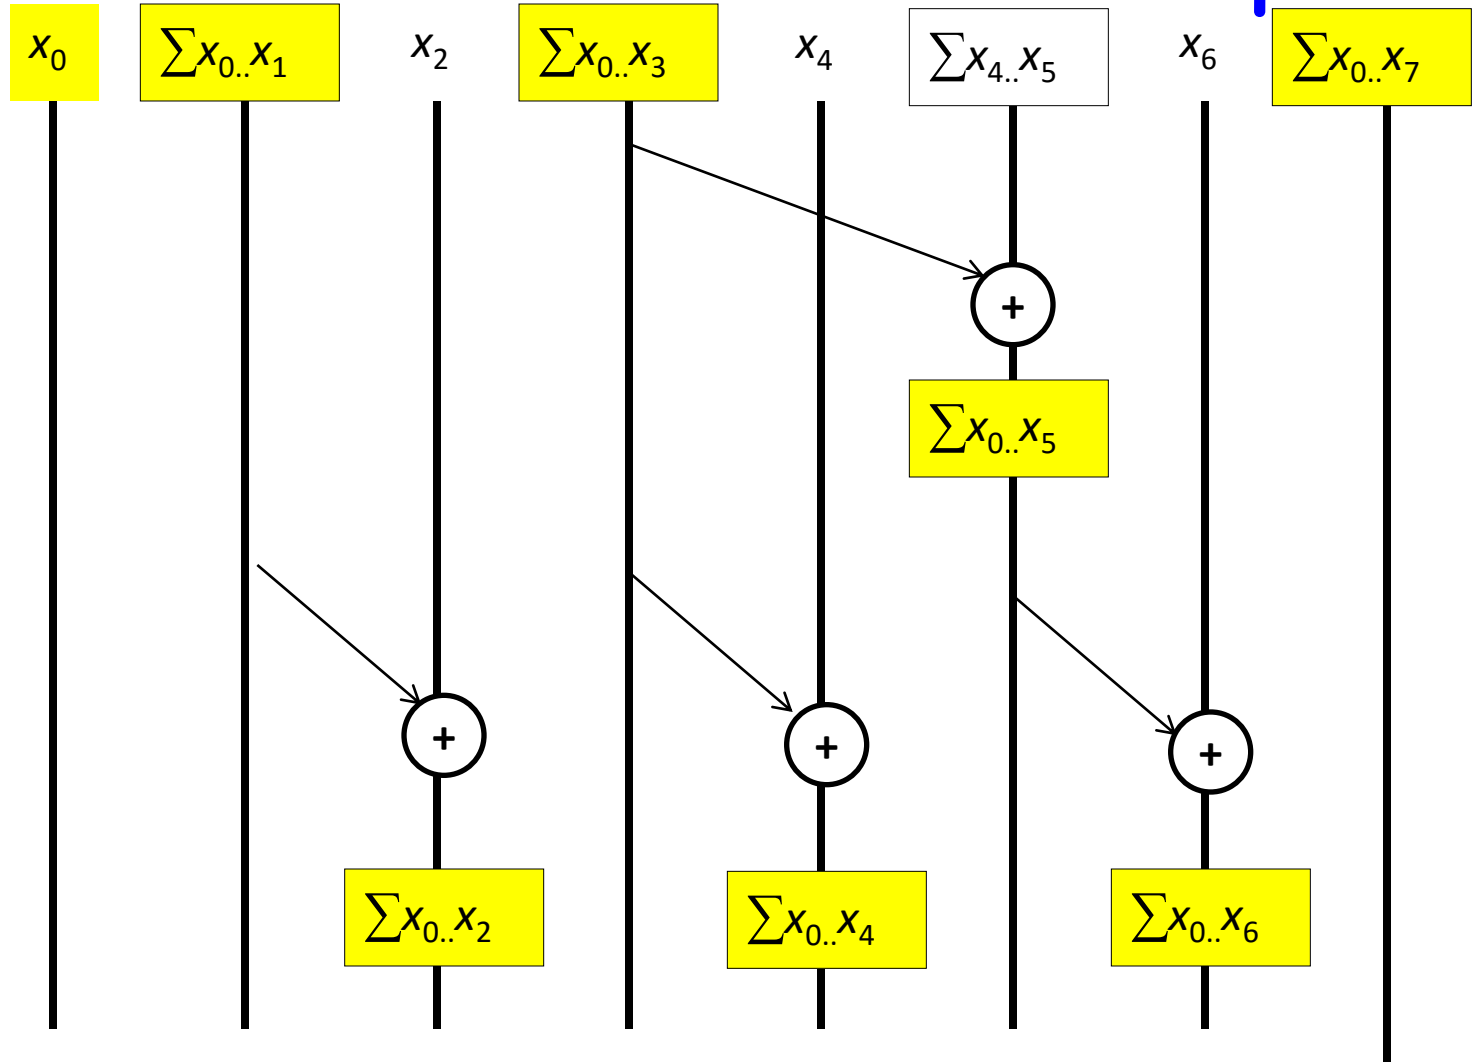
# Parallel Scan - Reduction Step




# Inclusive Post Scan Step



# Inclusive Post Scan Step



# Putting it Together



# Work Analysis

- The parallel Scan executes  $2 \cdot \log(n)$  parallel iterations
  - $\log(n)$  in reduction and  $\log(n)$  in post scan
  - The iterations do  $n/2, n/4, \dots, 1, 1, \dots, n/4, n/2$  adds
  - Total adds:  $2 \cdot (n-1) \rightarrow O(n)$  work
- The total number of adds is no more than twice of that done in the efficient sequential algorithm
  - The benefit of parallelism can easily overcome the 2X work when there is sufficient hardware

# Error Handling in *CUDA*



```
__global__ void foo(int *ptr)
{
 *ptr = 7;
}

int main(void)
{
 foo<<<1,1>>>(0);
 return 0;
}
```

What will happen when you compile and execute this piece of code?

# Error Handling

- In a CUDA program, if we suspect an error has occurred during a kernel launch, then we **must explicitly check** for it after the kernel has executed.
- CUDA runtime will respond to questions ... But won't talk without asked!

`cudaError_t cudaGetLastError(void);`

- Called by the host
- returns a value encoding the kind of the last error it has encountered
- check for the error only after we're sure a kernel has finished executing → don't forget kernel calls are async!
  - What will you do?

```
#include <stdio.h>
#include <stdlib.h>
```

```
__global__ void foo(int *ptr)
{
 *ptr = 7;
}
```

```
int main(void)
{
 foo<<<1,1>>>(0);
```

```
// make the host block until the device is finished with foo
cudaThreadSynchronize();
```

```
// check for error
cudaError_t error = cudaGetLastError();
if(error != cudaSuccess)
{
 // print the CUDA error message and exit
 printf("CUDA error: %s\n", cudaGetErrorString(error));
 exit(-1);
}
```

```
return 0;
}
```

```
$ nvcc crash.cu -o crash
```

```
$./crash
```

```
CUDA error: unspecified launch failure
```

# Same Technique with Synchronous Calls

```
cudaError_t error = cudaMalloc((void**)&ptr,
 10000000000000);
if(error != cudaSuccess)
{
 // print the CUDA error message and exit
 printf("CUDA error: %s\n",
 cudaGetErrorString(error));
 exit(-1);
}
```

The output will be:

**CUDA error: out of memory**

# Rules of Thumb

- Do not use `cudaThreadSynchronize()` a lot in your code because it has a large performance penalty.
- You can enable it during debugging and disable it otherwise.

**#ifdef DEBUG**

```
cudaThreadSynchronize();
cudaError_t error = cudaGetLastError();
if(error != cudaSuccess)
{
 printf("CUDA error at %s:%i: %s\n", filename, line_number, cudaGetErrorString(error));
 exit(-1);
}
```

**#endif**

If debugging, compile with:  
\$ nvcc **-DDEBUG** mycode.cu

# Conclusions

- Performance is related to how you keep the GPU and its memory busy → does not necessarily mean higher occupancy.
- We looked at some of the common parallel patterns used in many GPU kernels. These are tools that you can use in your own kernels.